

Lecture 15

Multiprocessors



Why Multiprocessors?

- Application requirements
 - Uniprocessor speed improving fast
 - But there are things that need even more speed
- Microprocessors as the fastest CPUs
 - Collecting several much easier than redesigning one
- Complexity of current microprocessors
 - Do we have enough ideas to sustain 1.5X/yr?
 - Can we deliver such complexity on schedule?
- Slow (but steady) improvement in parallel software
 - (scientific apps, databases, OS)



Opportunities for Applications

■ Scientific computing

- **Nearly unlimited demand** (Grand challenge)


App	Perf(GFLOPS)	Memory (GB)
72 hour weather	3	1
Pharmaceutical design	100	10
Global Change,Genome	1000	1000

■ Successes in some real industries:

- Petroleum reservoir modeling
- Automotive: crash simulation, drag analysis, engine
- Aeronautics: airflow analysis, engine, structural mechanics
- Pharmaceuticals: molecular modeling
- Entertainment: full length movies (“Toy Story”)

■ Commercial application

- Transaction processing, file servers, electronic CAD simulation,search engine
- Examples: IBM RS6000, Tandem (Compaq) Himilaya



Parallel Processing

- Multiple processors working cooperatively on problems: *not* the same as multiprogramming
- Goals/Motivation
 - Performance: limits of uniprocessors
ILP (branch prediction, RAW dependencies, memory)
 - Cost Efficiency: build big systems with commodity parts
(uniprocessors)
 - Scalability: just add more processors to get more performance
 - Fault tolerance: One processor fails you still can continue processing

Parallel Computers

- **Definition:** “A parallel computer is a collection of processing elements that cooperate and communicate to solve large problems fast.”

Almasi and Gottlieb, *Highly Parallel Computing*, 1989

- Questions about parallel computers:
 - How large a collection?
 - How powerful are processing elements?
 - How do they cooperate and communicate?
 - How are data transmitted?
 - What type of interconnection?
 - What are HW and SW primitives for programmer?
 - Does it translate into performance?

Popular Flynn Taxonomy

■ Flynn's Taxonomy of Parallel Machines

- How many Instruction streams?
- How many Data streams?

■ SISD (Single Instruction Single Data) ----Uniprocessors

■ MISD (Multiple Instruction Single Data) ---- Not used much

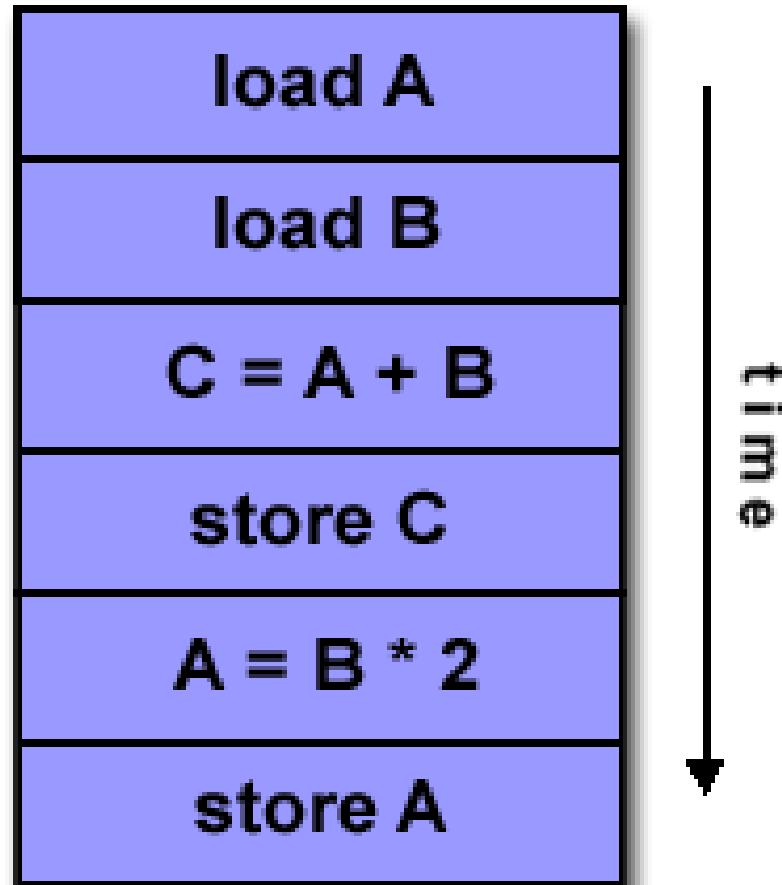
■ SIMD (Single Instruction Multiple Data)

- Each “processor” works on its own data, but execute the same instr.
- Examples: connection machine 2: 65535↑ 1bit processors;
Illiad IV: 64↑ 64bit processors;
- Ad: Simple programming model; Low overhead; Flexibility;

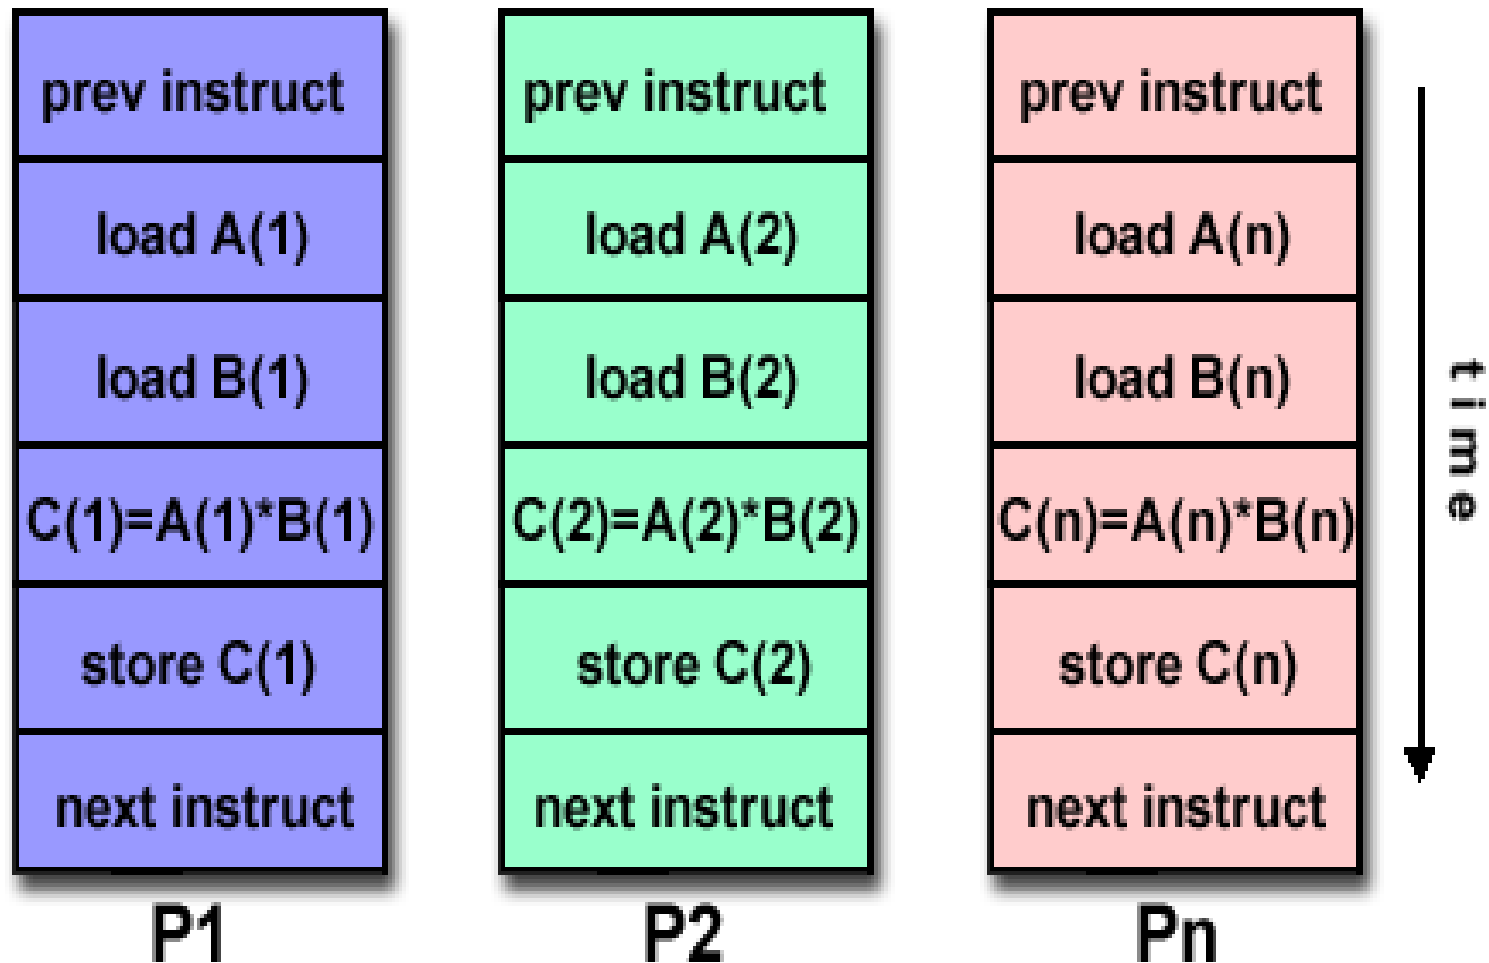
■ MIMD (Multiple Instruction Multiple Data)

- Each processor executes its own instr. and operates on its own data
- Examples: Sun Enterprise 5000, Cray T3D, SGI Origin
- Ad: Flexible; Use off-the-shelf microprocessors;
- *Not superscalar, each node is superscalar, lessons will apply to multi-core*

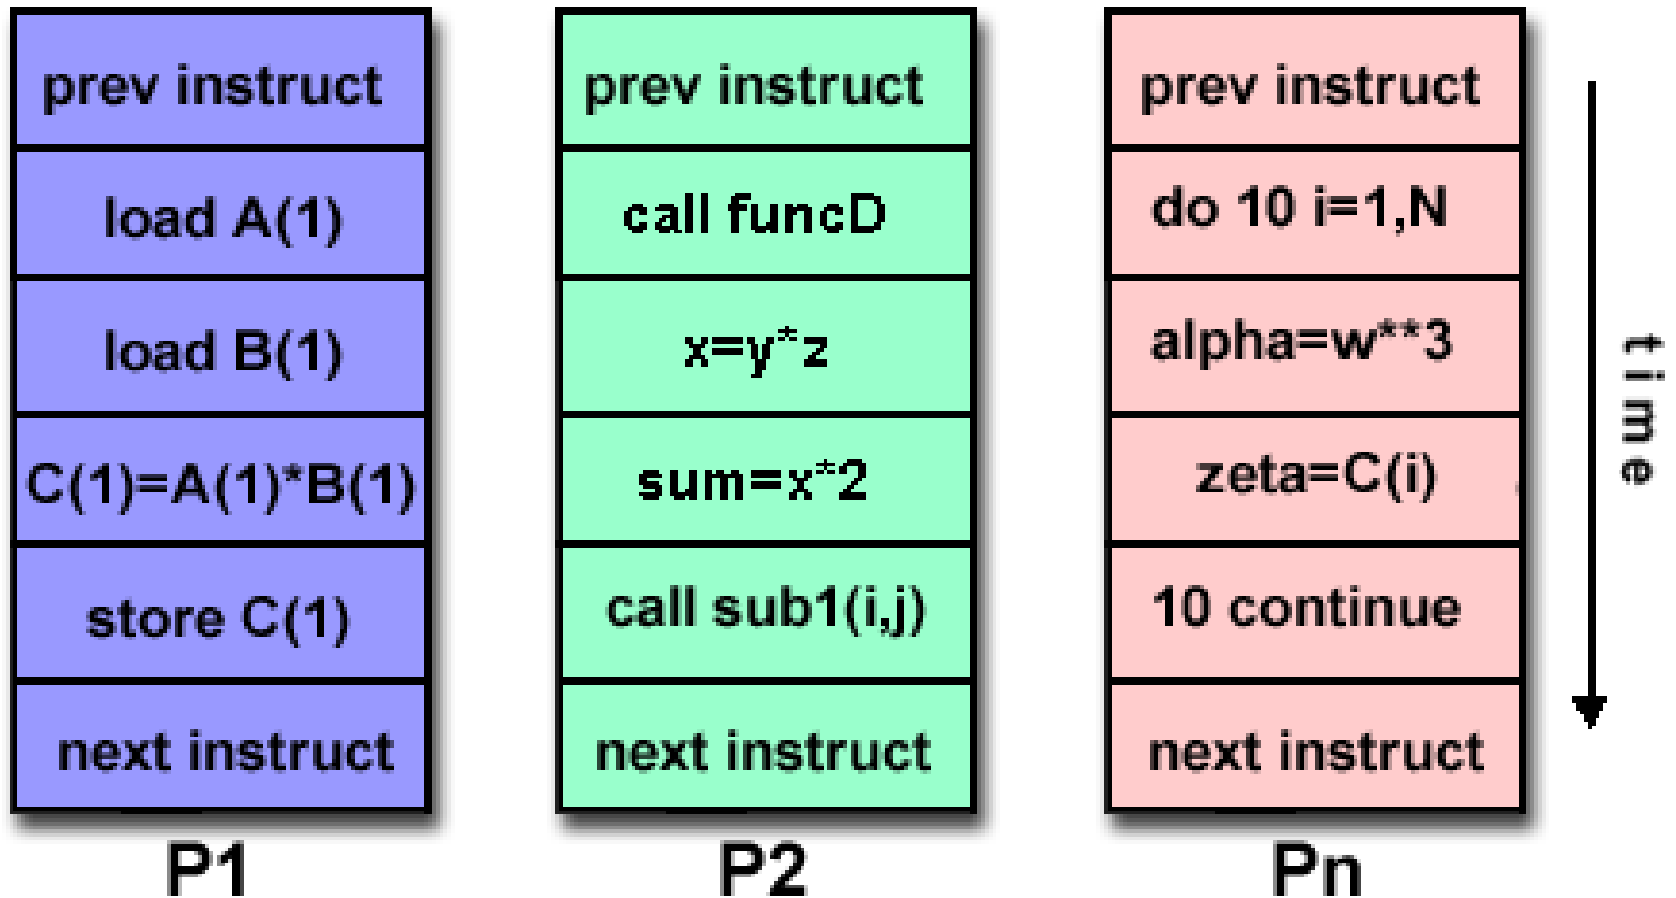
Program execute on SISD



Program execute on SIMD



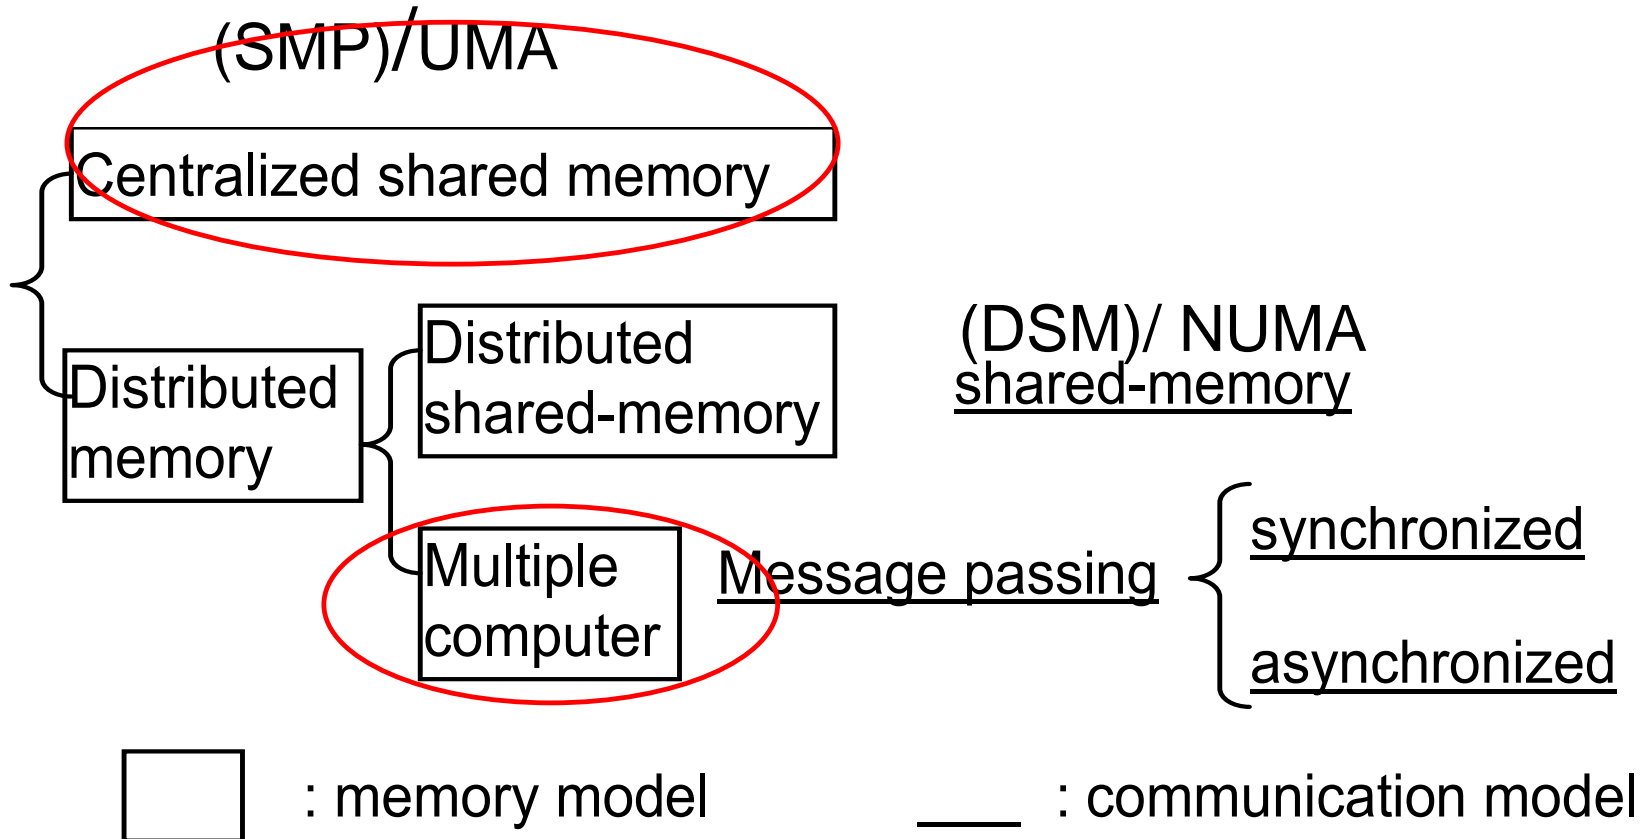
Program execute on MIMD




Catalogue the Parallel (MIMD) processors

- Center on organization of main memory
 - Shared vs. Distributed
- Appearance of memory to hardware
 - Q1: Memory access latency uniform?
 - Shared : yes, doesn't matter where data goes
 - Distributed: no, makes a big difference
- Appearance of memory to software
 - Q2: Can processors communicate directly via memory?
 - Shared (shared memory): yes, communicate via load/store
 - Distributed (message passing): no, communicate via messages
- Dimensions are orthogonal
 - e.g. DSM: (physically) distributed, (logically) shared memory

Dimensions are orthogonal




UMA vs. NUMA: Why it matters



- Ideal model:
 - Perfect (single-cycle) memory latency
 - Perfect (infinite) memory bandwidth
- Real systems:
 - Latencies are long and grow with system size
 - Bandwidth is limited
 - Add memory banks, interconnect to hook up (latency goes up)

UMA vs. NUMA



■ UMA: uniform memory access

- From p0 same latency to m0-m3
- Data placement doesn't matter
- Latency worse as system scales
- Interconnect contention restricts bandwidth
- Small multiprocessors only

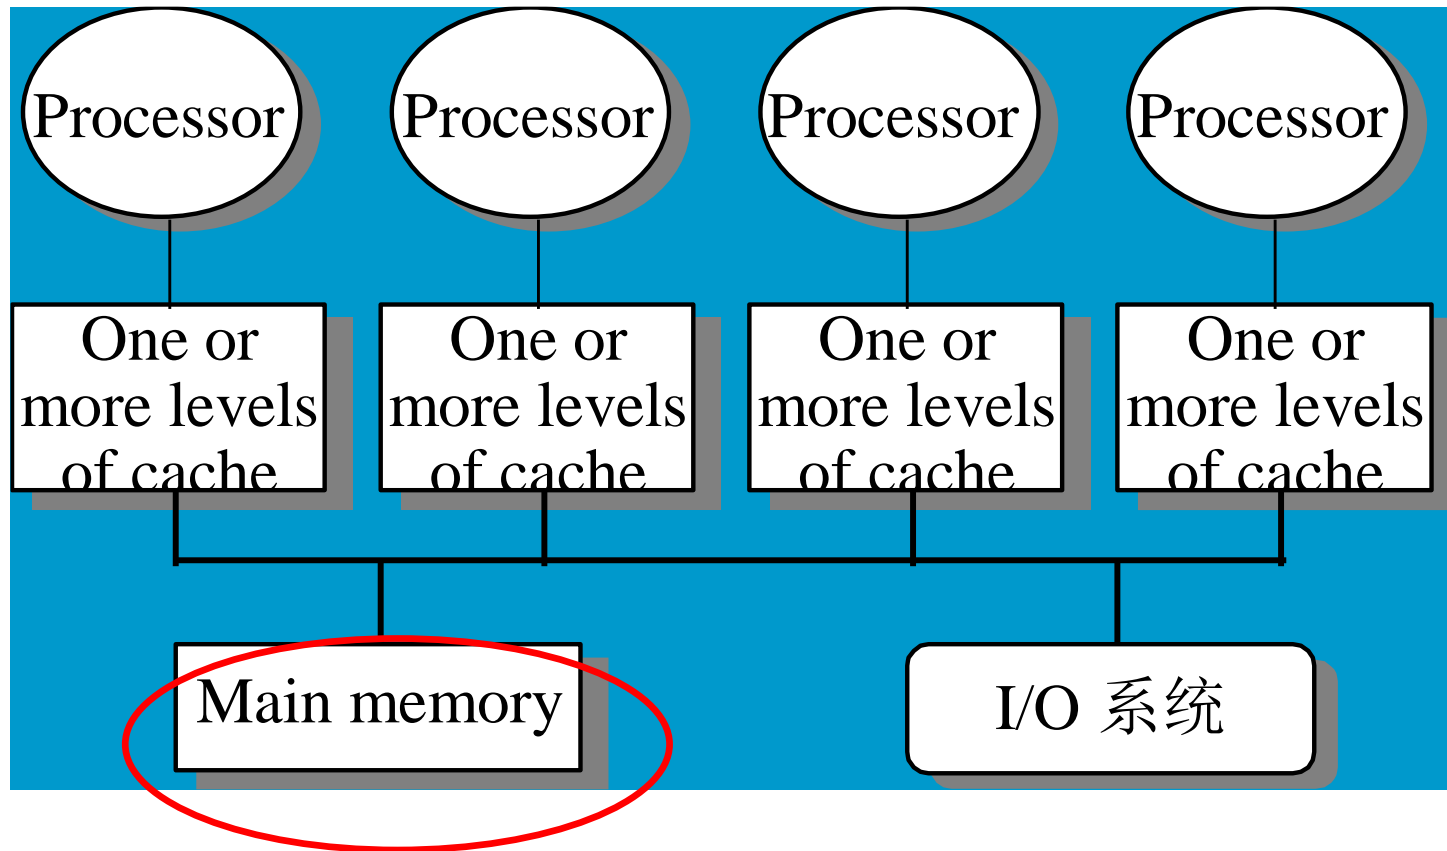
■ NUMA: non-uniform memory access

- From p0 faster to m0 than m1-m3
- Low latency to local memory helps performance
- Data placement important (software!)
- Less contention => more scalable
- Large multiprocessor systems

Major MIMD Styles

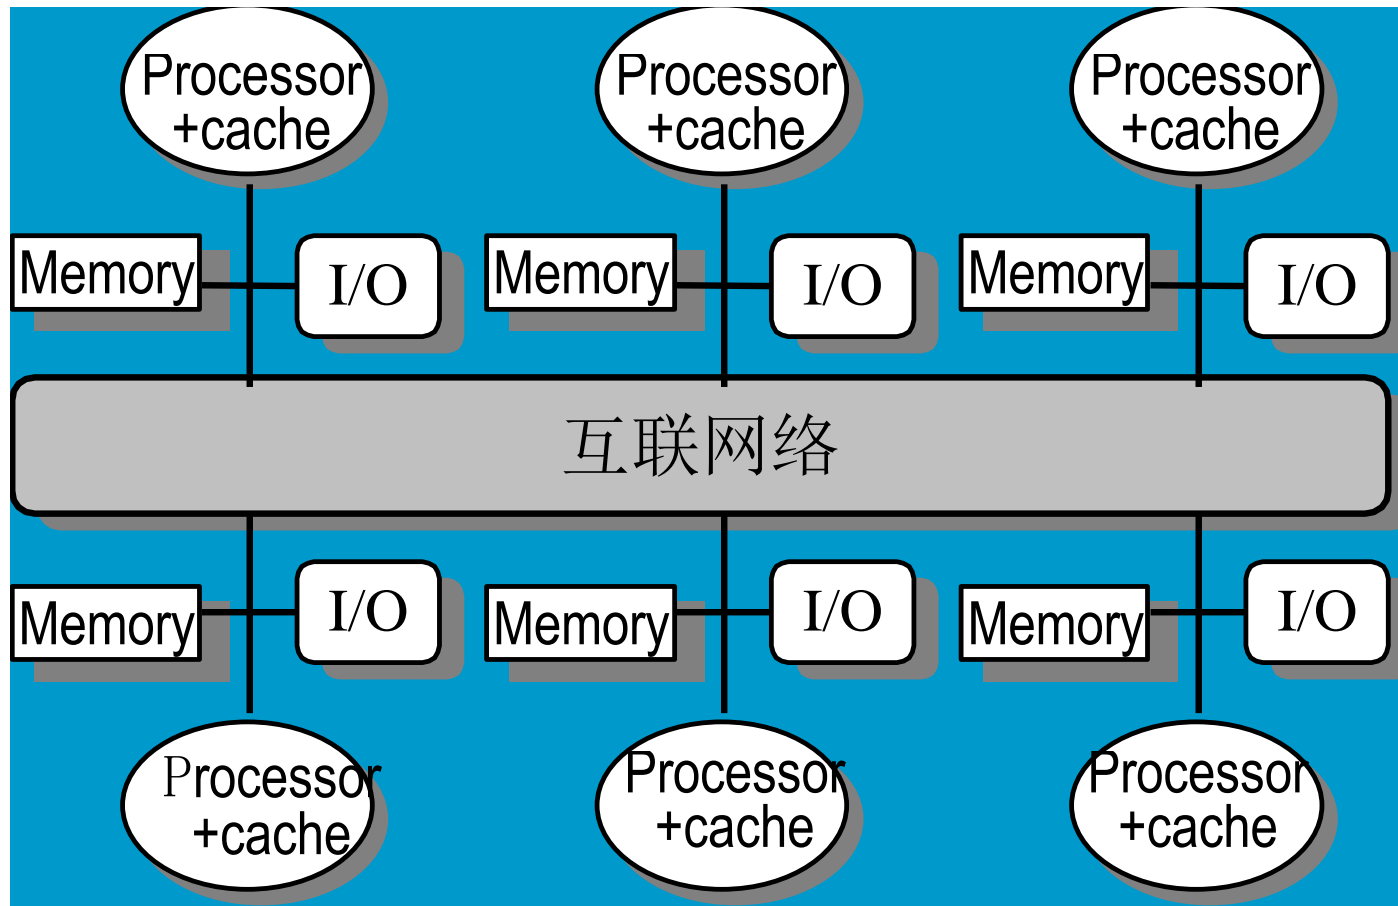
- **Centralized shared memory** ("Uniform Memory Access" time or "Shared Memory Processor")
- **Decentralized memory** (memory module with CPU)
 - get more memory bandwidth, lower memory latency
 - Drawback: Longer communication latency
 - Drawback: Software model more complex

Structure of centralized shared-memory multiprocessor




< 100 processor nodes in 2006, normal few dozen

Structure of distributed-memory multiprocessor




Comparison in graph


- Uniprocessor




- Pipelined




- Superscalar




- VLIW



- SMP(Symmetric) or CSM(Centralized)



- Distributed



Distributed-memory multiprocessor(1)

■ Distributed shared memory

(DSM or scalable shared memory)

- logical uniform address space but physical distributed memory, so any one of the processors can access any one of the memories.
- Shared memory means sharing the address space, which is different from centralized shared memory.

Distributed-memory multiprocessor(2)

■ multiple computers

- Address space consists of **multiple private (separate) address spaces**. A processor can **NOT** access the remote memory node directly.
- Every node (processor-memory pair) is a independent computer.
- NOW(Network of Workstation)is consisted of multiple node(PC or workstation) connected by LAN
- PC cluster

Parallel Architecture

- Parallel Architecture extends traditional computer architecture with a **communication architecture**
 - abstractions (HW/SW interface)
 - organizational structure to realize abstraction efficiently

Parallel Framework

■ Programming Model:

- **Multiprogramming** : lots of jobs, no communication
- **Shared address space**: communicate via memory
- **Message passing**: send and receive messages
- **Data Parallel**: several agents operate on several data sets simultaneously and then exchange information globally and simultaneously (shared or message passing)

■ Communication Abstraction:

- **Shared address space**: e.g., load, store, atomic swap
- **Message passing**: e.g., send, receive library calls

Shared Address Model-1

- Each **processor** can name every **physical location** in the machine
- Each **process** can name all data that shares with other processes
- Data transfer via **load and store**
- Data size: byte, word, ... or cache blocks
- Uses virtual memory to map virtual space to local or remote physical space
- Memory hierarchy model applies: now communication moves data to local processor cache (as load moves data from memory to cache)
 - Latency, Bandwidth, scalability when communicate?
- For distributed memory architecture, a layer (software or hardware) is generally added to allow transparent address mapping

Shared Address Model-2

- Significant research has been conducted to make the translation transparent and scalable for many node
- **Handling data consistency and protection is typically a challenge**
- For multi-computer systems, address mapping has to be performed by software modules, typically added as part of the operating system
- Latency depends on the underlined hardware architecture (bus bandwidth, memory access time and support for address translation)
- **Scalability is limited** given that the communication model is so tightly coupled with process address space*

Message Passing Model-1

- Whole computers (CPU, memory, I/O devices) communicate as **explicit I/O operations**
 - Essentially NUMA but integrated at I/O devices vs. memory system
- **Send** specifies local buffer + receiving process on remote computer
- **Receive** specifies sending process on remote computer + local buffer to place data
 - Usually send includes process tag and receive has rule on tag: match 1, match any
 - **Synch**: when send completes, when buffer free, when request accepted, receive wait for send
- **Send + receive => memory-memory copy, where each supplies local address, AND does pairwise synchronization!**

Message Passing Model-2

■ History of message passing:

- Network topology important because could only send to immediate neighbor
- Typically synchronous, blocking send & receive
- Later DMA with non-blocking sends, DMA for receive into buffer until processor does receive, and then data is transferred to local memory
- Later SW libraries to allow arbitrary communication

■ Example:

- IBM SP-2, RS6000 workstations in racks
- Network Interface Card has Intel 960
- 8X8 Crossbar switch as communication building block
- 40 MByte/sec per link

Shared Memory vs. Message Passing

■ Shared Memory (multiprocessors)

- One shared address space
- Processors **use conventional load/stores to access shared data**
- Communication can be complex/dynamic
- **Simpler programming model** (compatible with uniprocessors)
- Hardware controlled caching is useful to reduce latency contention
- Has **drawbacks**
 - **Synchronization** (discussed later)
 - **More complex hardware** needed

Shared Memory vs. Message Passing

- MIMD (appearance of memory to software)
- Message Passing (multicomputers)
 - Each processor has its own address space
 - Processors send and receive messages to and from each other
 - Communication patterns explicit and precise
 - Explicit messaging forces programmer to optimize this
 - Used for scientific codes (explicit communication)
 - Message passing systems: PVM, MPI, OpenMP
 - Simple Hardware
 - Difficult programming Model

Communication Models

■ Shared Memory

- Processors communicate with shared address space
- Easy on small-scale machines
- Advantages:
 - Model of choice for uniprocessors, small-scale MPs
 - Ease of programming
 - Lower latency
 - Easier to use hardware controlled caching

■ Message passing

- Processors have private memories, communicate via messages
- Advantages:
 - Less hardware, easier to design
 - Focuses attention on costly **non-local** operations

■ Can support either SW model on either HW base

Parallel Systems (80s and 90s)

Machine	Communication	Interconnect	#cpus	Remote latency (us)
SPARCcenter	Shared memory	Bus	≤ 20	1
SGI Challenge	Shared memory	Bus	≤ 32	1
CRAY T3D	Shared memory	3D Torus	64-1024	1
Convex SPP	Shared memory	X-bar/ring	8-64	2
KSR-1	Shared memory	Bus/ring	32	2-6
TMC CM-5	Messages	Fat tree	64-1024	10
Intel Paragon	Messages	2D mesh	32-2048	10-30
IBM SP-2	Messages	Multistage	32-256	30-100

Multiprocessor Trends

■ Shared Memory

- Easier, more dynamic programming model
- Can do more to optimize the hardware

■ Small-to-medium size UMA systems (2-8 processors)

- Processor + memory + switch on single board (4x pentium)
- Single-chip multiprocessors (POWER4)
- Commodity parts soon – glueless MP systems

■ Larger NUMAs built from smaller UMAs

- Use commodity small UMAs with commodity interconnects (ethernet, myrinet)
- NUMA clusters

Fundamental Issues_1

- *Naming*: how to solve large problem fast
 - what data is shared
 - how it is addressed
 - what operations can access data
 - how processes refer to each other
- Choice of naming affects **code** produced by a compiler; via load where just remember address or keep track of processor number and local virtual address for message passing
- Choice of naming affects **replication of data**; via load in cache memory hierarchy or via SW replication and consistency

Fundamental Issues_2

- **Global physical address space:** any processor can generate, address and access it in a single operation
- **Global virtual address space:** if the address space of each process can be configured to contain all shared data of the parallel program
- memory can be anywhere: virtual address translation handles it
- **Segmented shared address space:** locations are named $\langle \text{process number, address} \rangle$
- uniformly for all processes of the parallel program

Fundamental Issues_3

- *Synchronization:*
- To cooperate, processes must coordinate
- Message passing is implicit coordination with transmission or arrival of data
- Shared address => additional operations to explicitly coordinate: e.g., write a flag, awaken a thread, interrupt a processor

Fundamental Issues_4

Latency and Bandwidth

1. Bandwidth

- Need high bandwidth in communication
- Match limits in network, memory, and processor
- Challenge is link speed of network interface vs. bisection bandwidth of network

2. Latency

- Affects performance, since processor may have to wait
- Affects ease of programming, since requires more thought to overlap communication and computation
- Overhead to communicate is a problem in many machines

3. Latency Hiding

- How can a mechanism help hide latency?
- Increases programming system burden
- Examples: overlap message send with computation, prefetch data, switch to other tasks

Challenge:

limited program parallism

$$\text{Speedup}_{\text{Overall}} = \frac{1}{\left((1 - \text{Fraction}_{\text{enhanced}}) + \frac{\text{Fraction}_{\text{enhanced}}}{\text{Speedup}_{\text{enhanced}}} \right)}$$

- Example: Achieve speedup of 80 x using 100 processors
 - $80 = 1 / [\text{Fracparallel} / 100 + 1 - \text{Fracparallel}]$
 - $\text{Frac parallel} = 0.9975$
 \Rightarrow only 0.25% of the work can be serial!

Can parallized part be imporved to 100 times considering the overheads ?

Challenge :

long communication latency

■ Given: (P203)

- 32-processor machine, with each processor cycle time=0.5ns (2GHz);
- remote reference time= 200ns;
- all references except those involving communication hit in local memory;
- **base CPI=0.5** (all reference hit in the cache);
- Processors are stalled on a remote request.
- 0.2% of the instructions involve a remote reference.

The multiprocessor with all local references is $1.3/0.5 = 2.6$ times faster.

■ Answer:

- **CPI = Base CPI + Remote request rate \times Remote request cost**
=0.5 + 0.2% \times Remote request cost

The remote request cost is:

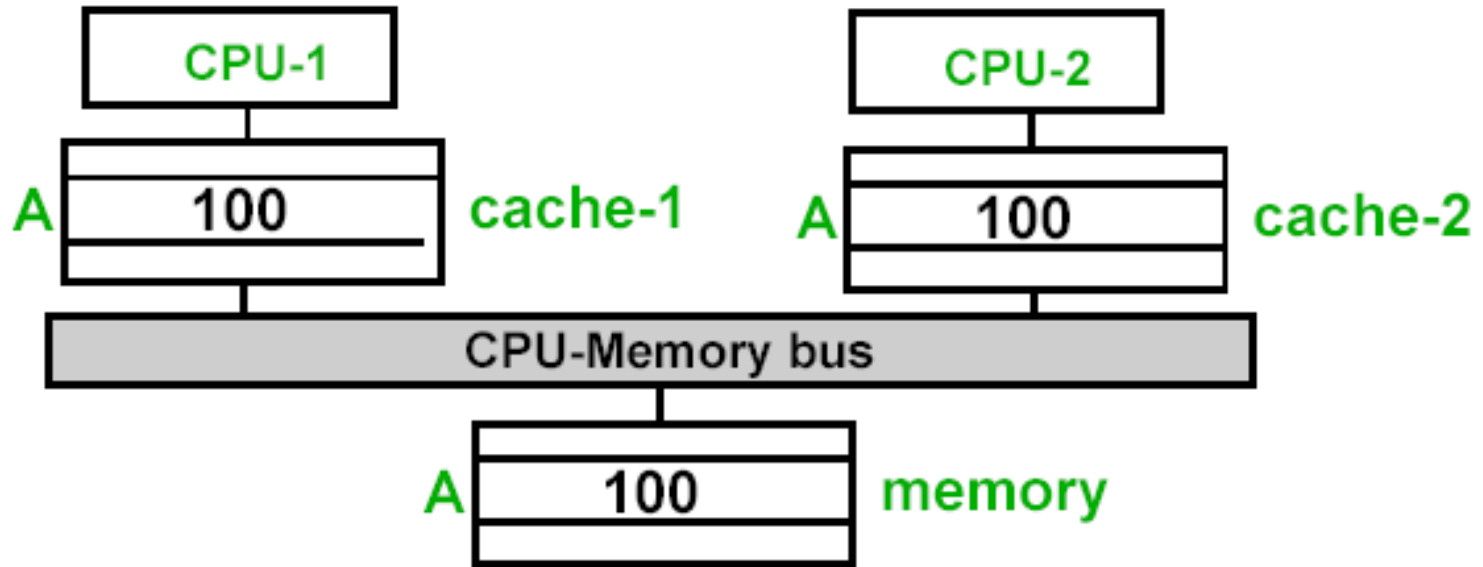
$$\frac{\text{Remote access cost}}{\text{Cycle time}} = \frac{200\text{ns}}{0.5\text{ns}} = 400\text{cycles}$$
$$\text{CPI} = 0.5 + 0.8 = 1.3$$

What can we do ?

- Limited program parallelism
 - New algorithm
- Long communication latency
 - HW: caching shared data to lower the remote access frequency
 - Problem: cache coherence, cache consistence
 - SW:
 - restructuring the data to make more accesses local
 - Synchronization
 - latency hiding techniques

Cache Coherence in Multiprocessor

Memory Consistency in SMPs



Suppose CPU-1 updates **A** to 200.

write-back: memory and cache-2 have stale values

write-through: cache-2 has a stale value

HW Coherence Protocols

■ Snooping Solution (Snoopy Bus):

- Send all requests for data to all processors
- Processors snoop to see if they have a copy and respond accordingly
- Requires broadcast, since caching information is at processors
- Works well with bus (natural broadcast medium)
- Dominates for small scale machines (most of the market)

■ Directory-Based Schemes (discuss later)

- Keep track of what is being shared in 1 centralized place (logically)
- Distributed memory => distributed directory for scalability (avoids bottlenecks)
- Send point-to-point requests to processors via network
- Scales better than Snooping
- Actually existed BEFORE Snooping-based schemes

Basic Snoopy Protocols

- Write Invalidate Protocol:
 - Multiple readers, single writer
 - Write to shared data: an invalidate is sent to all caches which snoop and invalidate any copies
 - Read Miss:
 - Write-through: memory is always up-to-date
 - Write-back: snoop in caches to find most recent copy
- Write Broadcast Protocol (typically write through):
 - Write to shared data: broadcast on bus, processors snoop, and update any copies
 - Read miss: memory is always up-to-date
- Write serialization: bus serializes requests!
 - Bus is single point of arbitration

EX: write back Cache, write invalidate

■ Mechanics

- Broadcast address of cache line to invalidate
- All processor snoop, then invalidate if in local cache
- policy can be used to service cache misses in write-back caches

Processor Activity	Bus activity	Contents of CPU A's cache	Contents of CPU B's cache	Contents of Memory Location X
				0
CPU A Reads X	Cache miss for X	0		0
CPU B Reads X	Cache miss for X	0	0	0
CPU A writes A 1 to X	Invalidation for X	1		0
CPU B Reads X	Cache miss for X	1	1	1

Ex: Write back Cache, update(Broadcast)

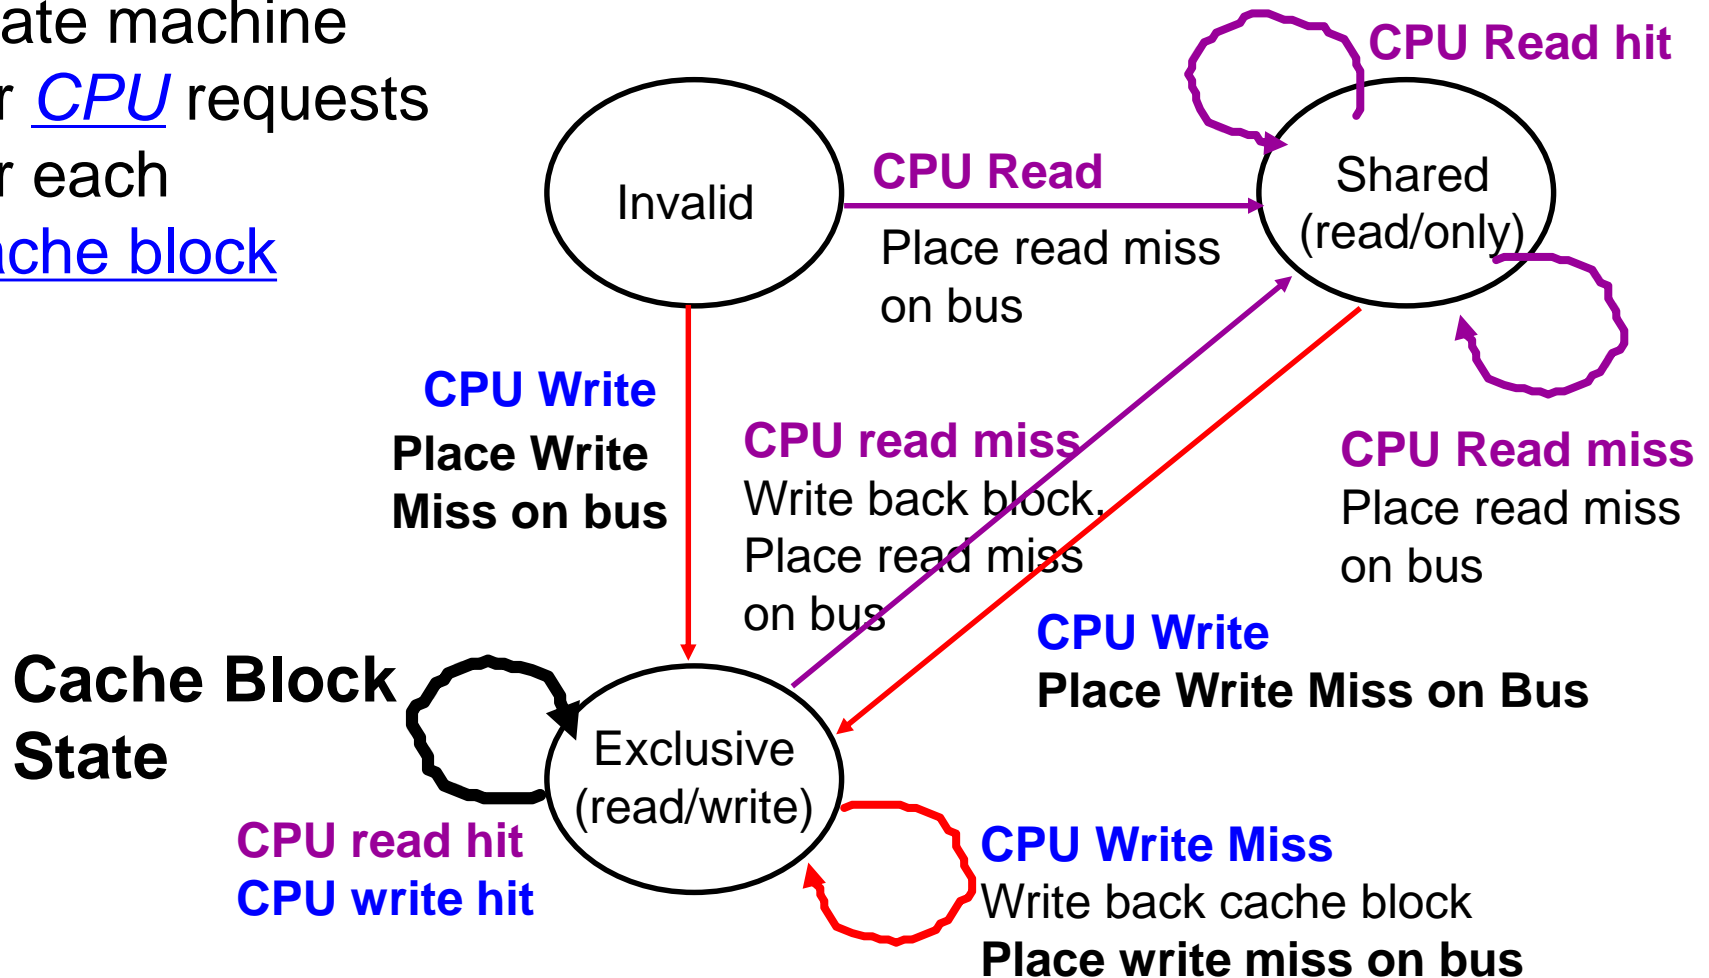
Processor Activity	Bus activity	Contents of CPU A's cache	Contents of CPU B's cache	Contents of Memory Location X
				0
CPU A Reads X	Cache miss for X	0		0
CPU B Reads X	Cache miss for X	0	0	0
CPU A writes A 1 to X	Write broadcast Of X	1	1	1
CPU B Reads X		1	1	1

Simple write-invalidate protocol

- Three states
 - Invalid, Shared, exclusive
- Events
 - CPU-R, CPU-W
 - BUS-R, BUS-W

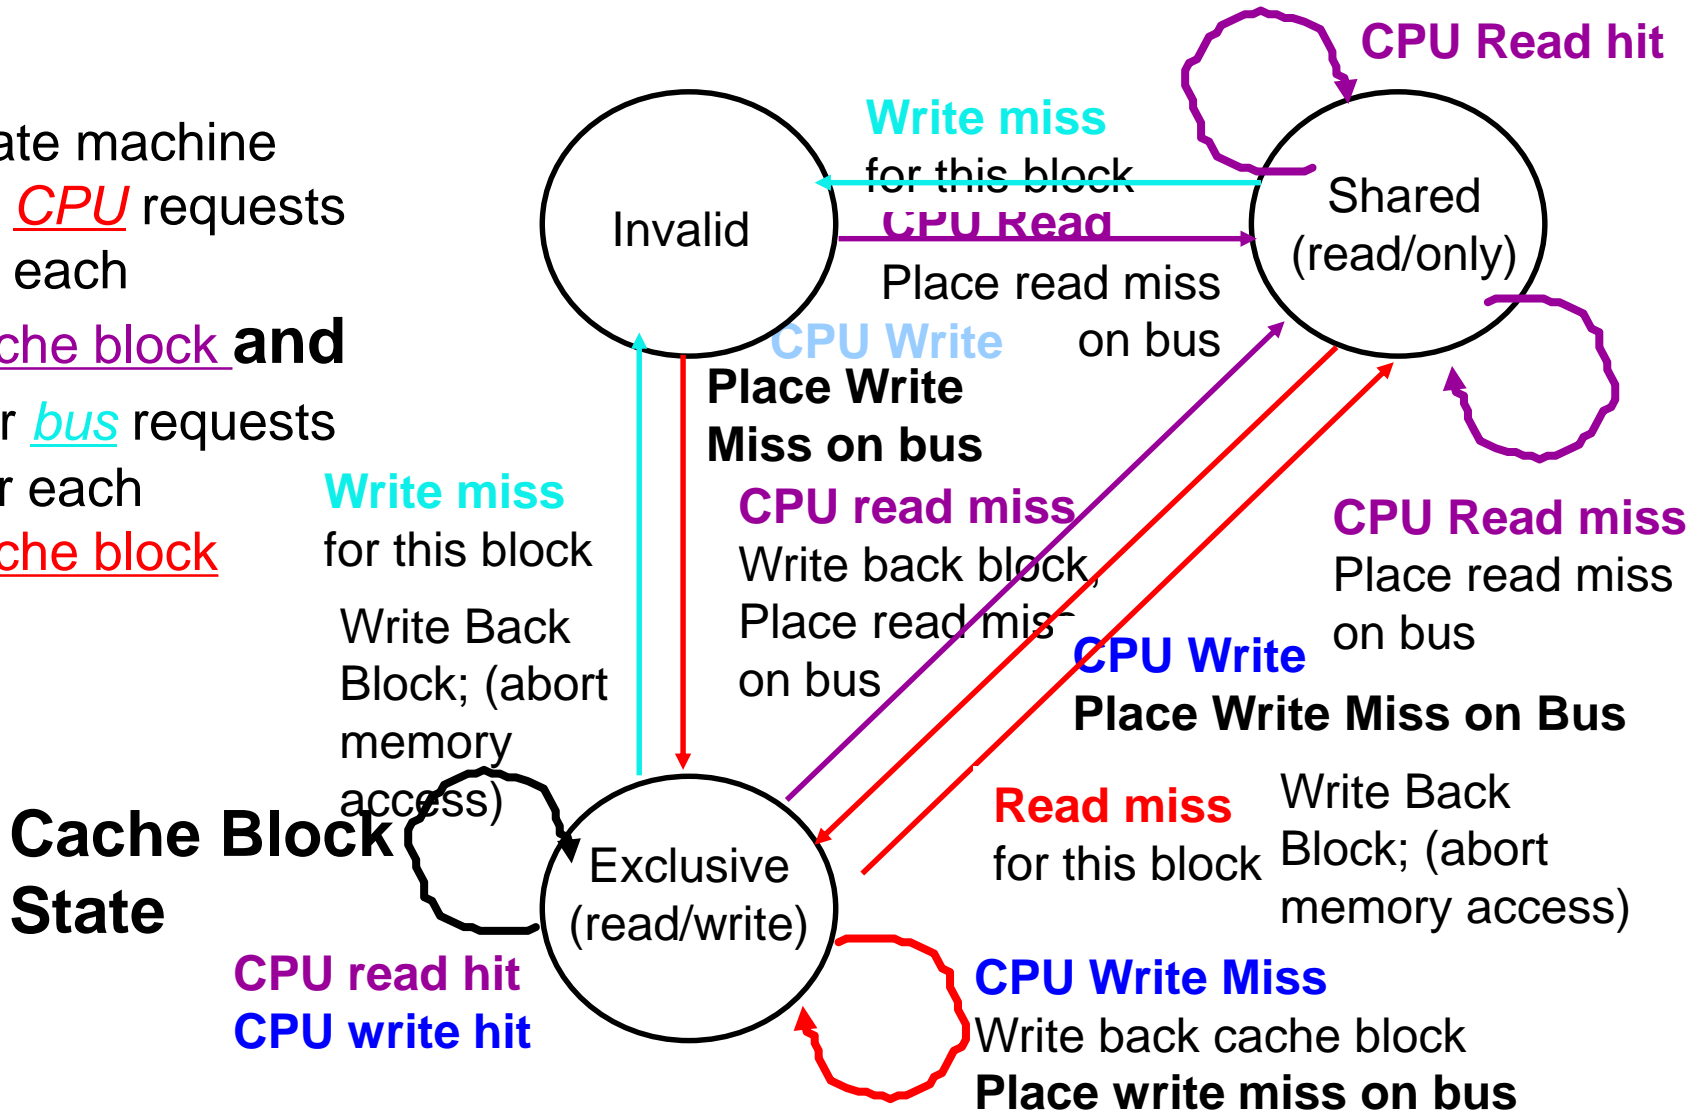
Snoopy-Cache State Machine-I

- State machine for CPU requests for each cache block

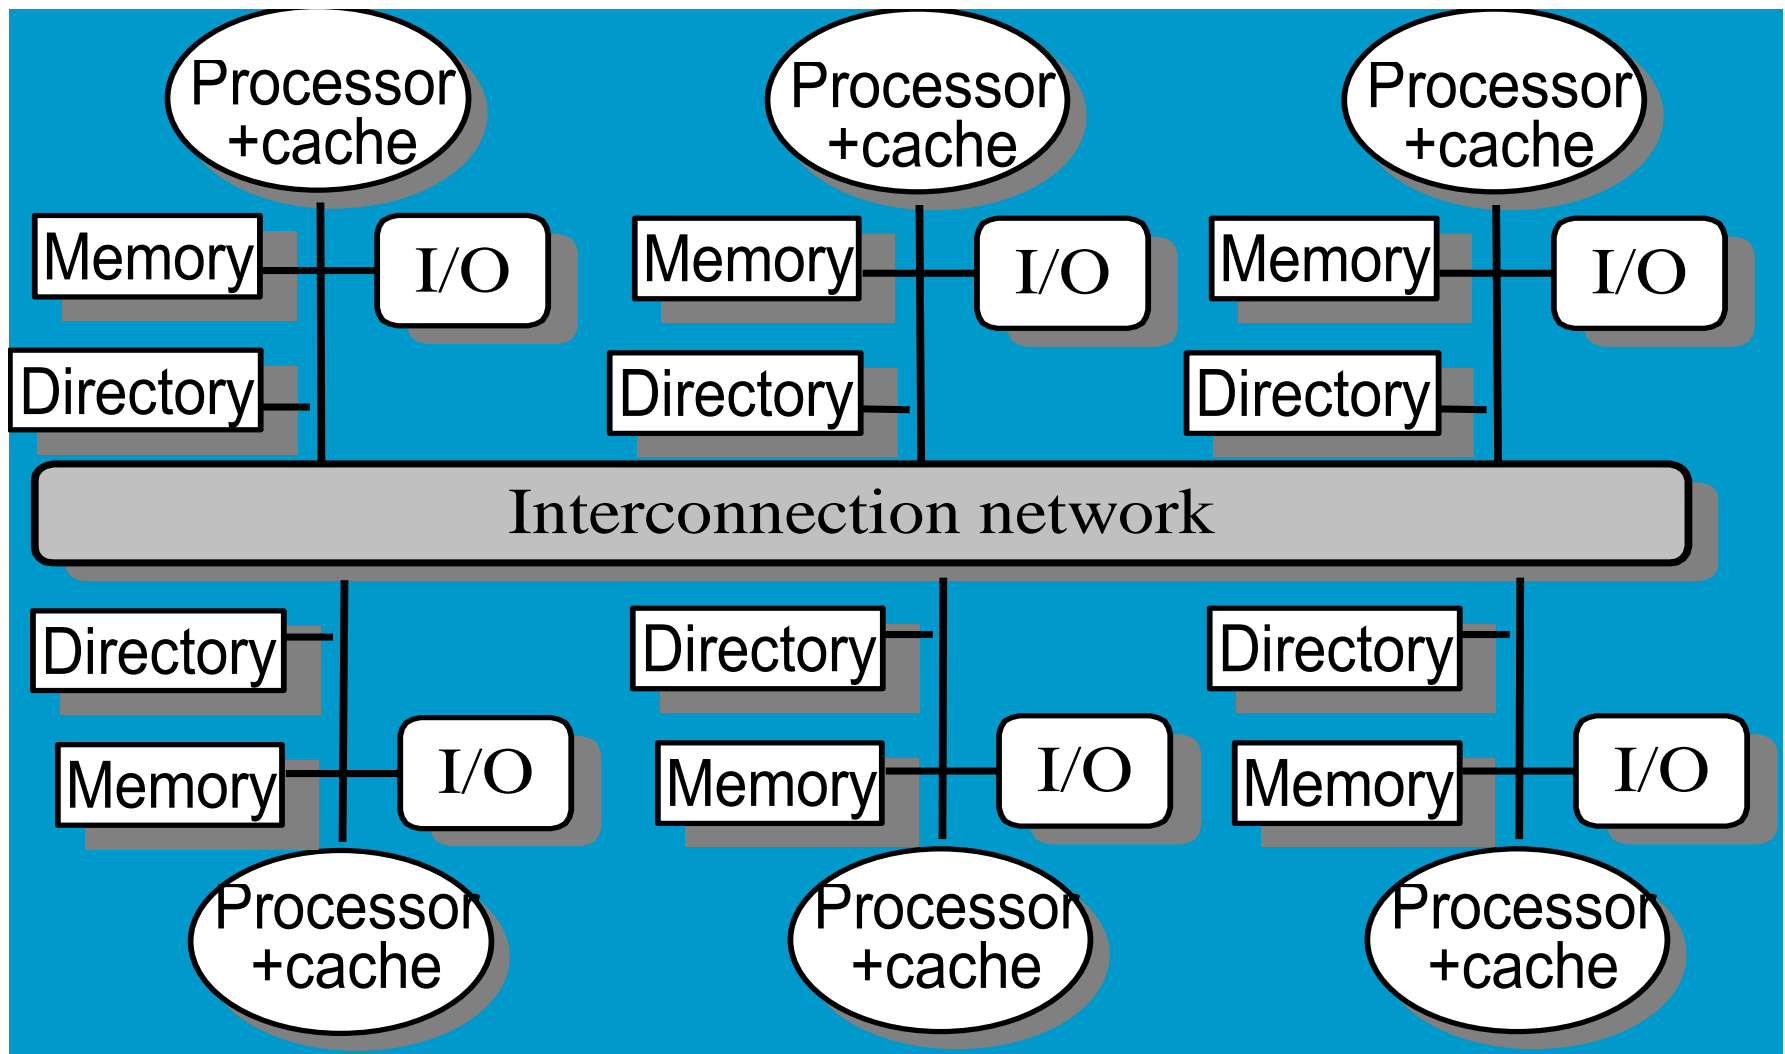


Snoopy-Cache State Machine-III

- State machine for CPU requests
 - for each cache block and **and**
 - for bus requests
 - for each **W**
 - cache block for



Distributed Directory MPs



Directory protocol

- **Directory:** track state of every block in memory, and change the state of block in cache according to directory.
- **Information in directory**
 - Status of Every block: **shared/uncached/exclusive**
 - **Which processors** have copies of the block: **bit vector**
 - Whether the block is dirty or clean
- Directory protocol can be implemented with a distributed memory
- Directory protocol can be applied to a centralized memory organized into banks

Directory protocol implementation

■ Block status

- Shared: ≥ 1 processors have data, memory up-to-date
- Uncached (no processor has it; not valid in any cache)
- Exclusive: 1 processor (**owner**) has data; memory out-of-date

■ Directory size = $f(\text{entry number} * \text{entry size})$

- Each memory block has an entry in directory / only keep the entries for cached blocks
- Every processor has one bit / Limited processor bits in bit vector

■ Directory can be distributed along with the memory to avoid becoming the bottleneck

■ Assumptions to Keep it simple:

- Writes to non-exclusive data \Rightarrow write miss
- Processor blocks until access completes
- Assume messages received and acted upon in order as sent

Directory Protocol

- No bus and don't want to broadcast:
 - Interconnect means no longer single arbitration point
 - all messages have explicit responses
- Terms: typically 3 processors involved
 - **Local node** where a request originates
 - **Home node** where the memory location of an address resides
 - **Remote node** has a copy of a cache block, whether exclusive or shared
- Example messages on next slide:
P = processor number, A = address

Message type	Source	Destination	Msg Content
Read miss	Local cache	Home directory	P, A
– Processor <i>P</i> reads data at address <i>A</i> ; make <i>P</i> a read sharer and arrange to send data back			
Write miss	Local cache	Home directory	P, A
– Processor <i>P</i> has a write miss at address <i>A</i> ; make <i>P</i> the exclusive owner and arrange to send data back			
Invalidate	Local cache	Home directory	A
– Request to send invalidates to all remote caches that are caching the block at address <i>A</i>			
Invalidate	Home directory	Remote caches	A
– Invalidate a shared copy at address <i>A</i> .			
Fetch	Home directory	Remote cache	A
– Fetch the block at address <i>A</i> and send it to its home directory			
Fetch/Invalidate	Home directory	Remote cache	A
– Fetch the block at address <i>A</i> and send it to its home directory; invalidate the block in the cache			
Data value reply	Home directory	Local cache	Data
– Return a data value from the home memory (read miss response)			
Data write-back	Remote cache	Home directory	A, Data
– Write-back a data value for address <i>A</i> (invalidate response)			

■ The End.