

Compiler Principle and Technology

Prof. Dongming LU
Apr. 22th, 2015

7. Runtime Environments

A series of horizontal lines in shades of red and white, extending across the width of the slide below the title.

Contents

Part One

7.1 Memory Organization During Program Execution

7.2 Fully Static Runtime Environments

7.3 Stack-Based Runtime Environments

Part Two

7.4 Dynamic Memory

7.5 Parameter Passing Mechanisms



- The previous chapters studied the phases of a compiler that **perform static analysis** of the source language
 - Scanning, parsing, and static semantic analysis
 - Depends only on the **properties of the source language**
- Now turn to the task of studying **how a compiler generates executable code**
 - **Additional analysis**, such as that performed by an optimizer
 - Some of this can be machine independent, but much of the task of code generation is **dependent on the details of the target machine**



- **Runtime Environment**

The structure of the target computer's **registers and memory** that serves to manage memory and maintain the information needed **to guide the execution process**

- **Three kinds of runtime environments**

(1) **Fully static** environment; FORTRAN77

(2) **Stack-Based** environment; C C++

(3) **Fully dynamic** environment; LISP



◆ Main issues discussed in the chapter:

³ For each environment, the language features and their properties

- (1) **Scoping and allocation** issues;
- (2) Nature of **procedure calls**;
- (3) **Parameter passing** mechanisms

◆ Focus on the **general structure** of the environment

◆ **Note:**

³ The compiler can only maintain an environment **only indirectly**

³ It must **generate code** to perform the necessary maintenance operations during program execution.




7.1 Memory Organization During Program Execution

A decorative graphic consisting of several horizontal lines of varying lengths and colors (dark red, light red, and white) extending from the left edge of the slide towards the right, positioned below the title.

The memory of a typical computer:

- A register area;
- Addressable Random access memory (RAM):



The code area is fixed prior to execution, and can be visualized as follows:

Entry pointer to procedure1→	Code for procedure 1
Entry pointer to procedure2→	Code for procedure 2
	...
Entry pointer to procedure n→	Code for procedure n

In particular, the entry point for each procedure and function is known at compile time.

- The global and/or static data of a program can be **fixed in memory prior to execution**
 - Data are allocated separately in a fixed area in a similar fashion to the code
 - In Fortran77, all data are in this class;
 - In Pascal, global variables are in this class;
 - In C, the external and static variables are in this class
- The **constants** are usually allocated memory **in the global/static area**
 - Const declarations of C and Pascal;
 - Literal values used in the code,
 - Such as “Hello%D\n” and Integer value 12345:
 - Printf(“Hello %d\n”,12345);

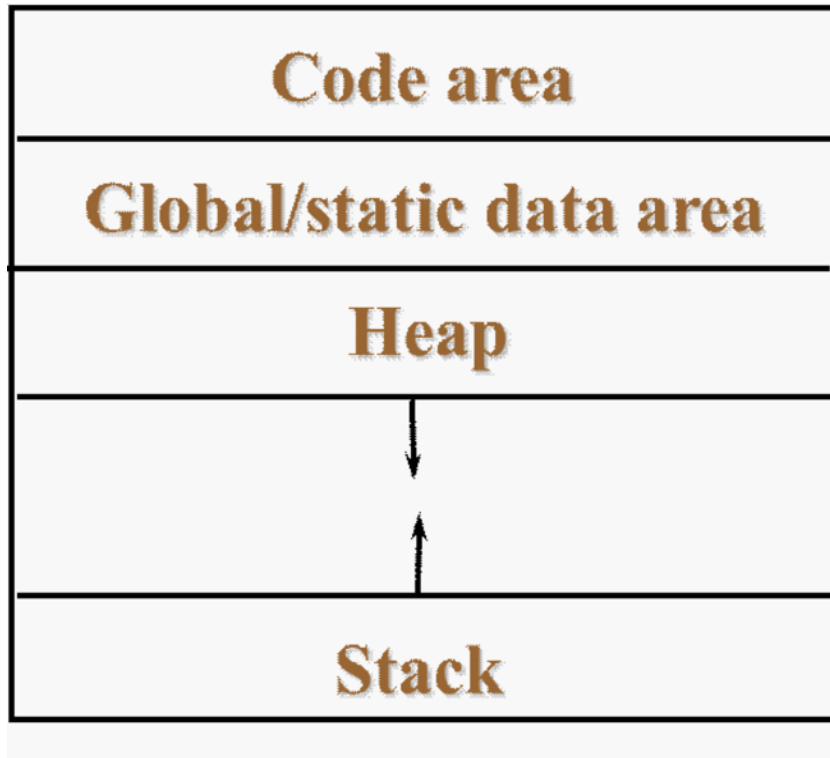


- ◆ The memory area used for **dynamic data** can be organized in many different ways
 - ³ **A stack area** used for data whose allocation occurs in FIFO fashion;
 - ³ **A heap area** used for dynamic allocation occurs not in FIFO fashion.

- ◆ The architecture of the target machine includes a **processor stack** for procedure calls and returns.
 - ³ A compiler will have to arrange for the **explicit allocation of the processor stack** in an appropriate place in memory.



The general organization of runtime storage:



Where, the arrows indicate the direction of **growth of the stack and heap.**

Procedure activation record

(An important unit of memory allocation)

Memory allocated for the local data of a procedure or function.

An activation record must contains the following sections:

Space for arguments (parameters)
Space for bookkeeping information, including return address
Space for local data
Space for local temporaries

Note: this picture only illustrates the general organization of an activation record.

- Some parts of an activation record have the **same size for all procedures**
 - Space for bookkeeping information
- Other parts of an activation record may **remain fixed for each individual procedure**
 - Space for arguments and local data
- Some parts of activation record may be **allocated automatically** on procedure calls:
 - Storing the return address
- Other parts of activation record may need to be **allocated explicitly** by instructions generated by the compiler:
 - Local temporary space
- **Depending on the language**, activation records may be allocated in different areas:
 - Fortran77 in the static area;
 - C and Pascal in the stack area; referred to as stack frames
 - LISP in the heap area.



- **Processor registers are also part of the structure of the runtime environment**
 - Registers may be used to store temporaries, local variables, or even global variables;
 - In newer RISC processor, keep entire static area and whole activation records;
 - Special-purpose registers to keep track of execution
 - PC program counter;
 - SP stack pointer;
 - FP frame pointer;
 - AP argument pointer



- **The sequence of operations when calling the functions: *calling sequence***
 - The allocation of memory for the activation record;
 - The computation and storing of the arguments;
 - The storing and setting of necessary registers to affect the call
- **The additional operations when a procedure or function returns: *return sequence (VS call)***
 - The placing of the return value where the caller can access it;
 - The readjustment of registers;
 - The possible releasing for activation record memory



◆ The important aspects of the design of the calling sequence:

(1) How to **divide the calling sequence** operations between the caller and callee

◆ At a minimum, the caller is responsible for computing the arguments and placing them in locations where they may be found by the callee

(2) To what extent to **rely on processor support for calls** rather than generating explicit code for each step of the calling sequence



7.2 Fully Static Runtime Environments

The entire program memory can be visualized as follows:

Code for main procedure	Code area
Code for procedure 1	
...	
Code for procedure n	
	Data area
Global data area	
Activation record of main procedure	
Activation record of procedure 1	
...	
Activation record of procedure n	

◆ All data are static, remaining fixed in memory for the duration of program execution

◆ For a language, such as FORTRAN77, **no pointer or dynamic allocation, no recursive procedure calling**

³ The global variables and all variables are allocated statically.

³ Each procedure has **only a single activation record**.

³ All variable, whether local or global, can be **accessed directly via fixed address**.



Advantages of Fully Static Runtime Environments

- Relative little overhead in terms **of bookkeeping information** to retain in each activation record;
- And **no extra information** about the environment needs to be kept in an activation record;
- The calling sequence is simple.
 - Each argument is computed and stored into its appropriate parameter location;
 - The return address is saved, and jump to the beginning of the code of the callee;
 - On return, a simple jump is made to the return address.



Example: A FORTRAN77 sample program

```
PROGRAM TEST
COMMON  MAXSIZE
INTEGER MAXSIZE
REAL    TABLE(10),TEMP
MAXSIZE = 10

READ *, TABLE(1),TABLE(2),TABLE(3)
CALL QUADMEAN(TABLE,3,TEMP)
PRINT *, TEMP
END
```

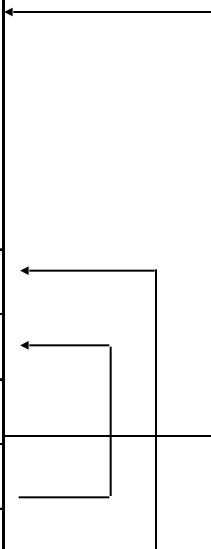


```
SUBROUTINE      QUADMEAN(A, SIZE,QMEAN)
COMMON  MAXSIZE
INTEGER  MAXSIZE,SIZE
REAL A(SIZE),QMEAN,TEMP
INTEGER K

TEMP=0.0
IF ((SIZE .GT. MAXSIZE) .OR. (SIZE .LT. 1) GOTO 99
DO 10 K=1,SIZE
TEMP=TEMP+A(K)*A(K)
10 CONTINUE
99 QMEAN = SQRT(TEMP/SIZE)
RETURN
END
```



A runtime environment for the program above.

Global area	MAXSIZE	
Activation record of main procedure	TABLE (1)	
	(2)	
	...	
	(10)	
Activation record of procedure QUADMEAN	TEMP	
	3	
	A	
	SIZE	
	QMEAN	
	Return address	
	TEMP	
	K	
	Unnamed location	

Note: The unnamed location is used to store temporary value during the computation of arithmetic expression.

limitations of Fully Static Runtime Environments

- Recursive calls are not allowed.
- The data object's size and their location in memory is decided when compiling.
- Dynamical memory allocation is not allowed.



7.3 Stack-Based Runtime Environments



- For a language, **in which**
 - Recursive calls are allowed;
 - Local variables are newly allocated at each call;
 - Activation records cannot be allocated statically
- Activation records must be **allocated in a stack-based fashion**
 - The stack of activation records grows and shrinks with the chain of calls in the executing program.
 - **Each procedure** may have several **different activation records on the call stack at one time**, each representing a distinct call.
 - **More complex strategy for bookkeeping and variable access**, which depends heavily on the properties of the language being compiled.



7.3.1 Stack-Based Environments Without Local Procedures



- All properties are global (such as C language), the stack-based environment requires two things:
 - (1) **Frame pointer or fp**, a pointer to the current activation record to allow access to local variable; **Control link or dynamic link**, a point to a record of the immediately preceding activation
 - (2) **Stack pointer or sp**, a point to the last location allocated on the call stack



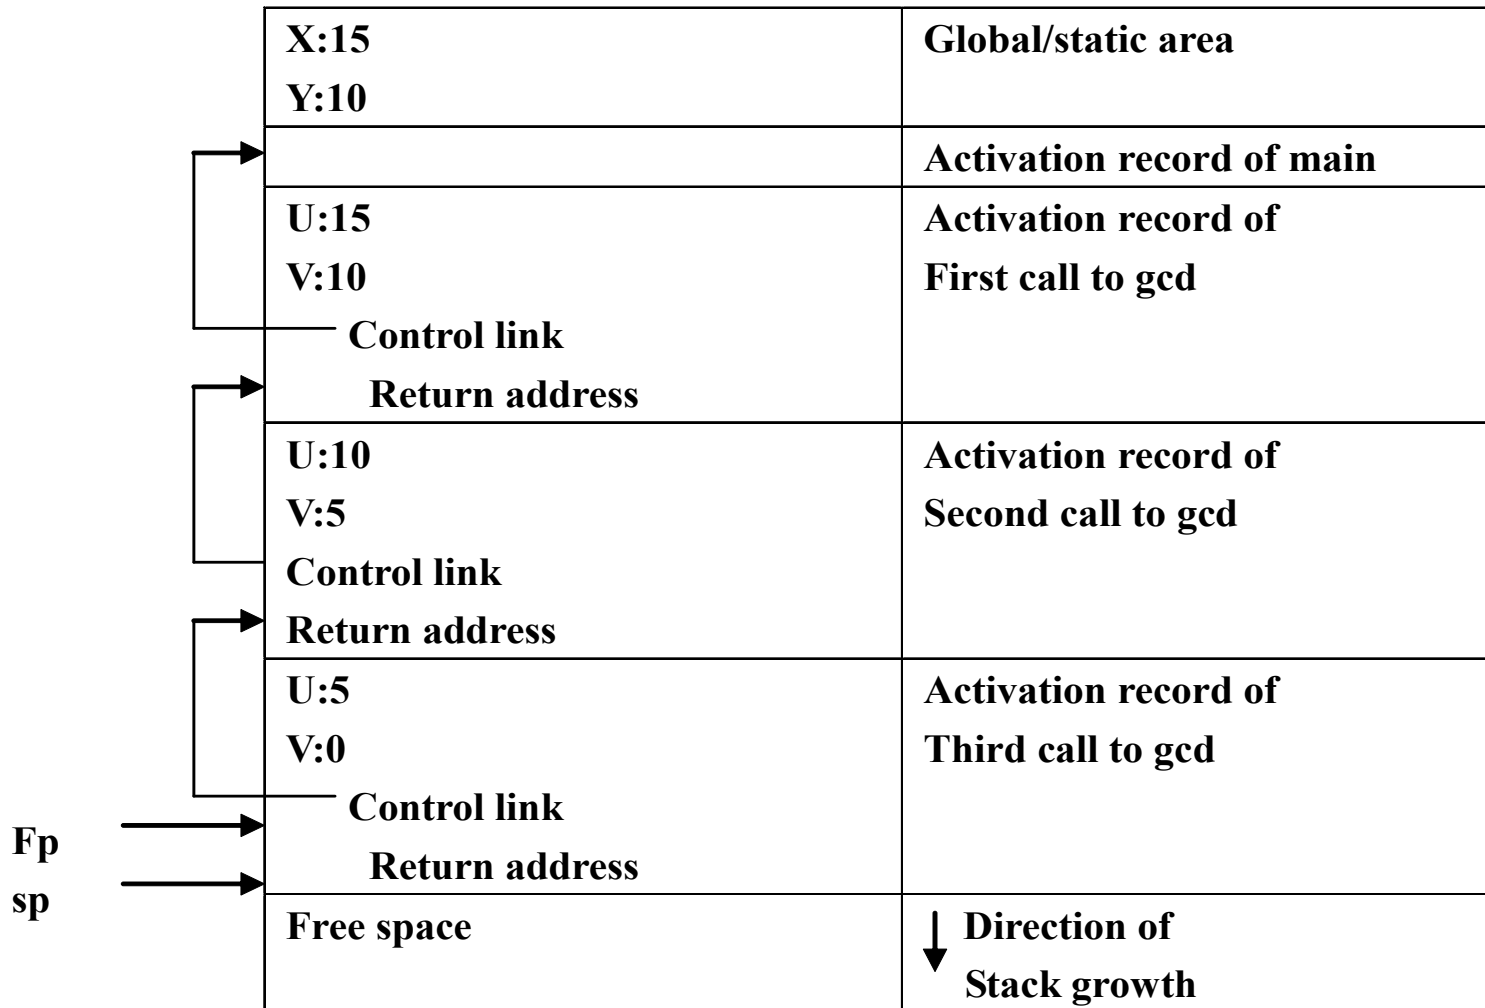
◆ **Example:** The simple recursive implementation of Euclid's algorithm to compute the greatest common divisor of two non-negative integer

```
# include <stdio.h>
int x,y;
int gcd(int u,int v)
{ if (v==0)
    return u;
    else return gcd(v,u%v);
}
main()
{ scanf("%d%d",&x,&y);
  printf("%d\n",gcd(x,y));
  return 0;
}
```

Suppose the user inputs the values 15 and 10 to this program.



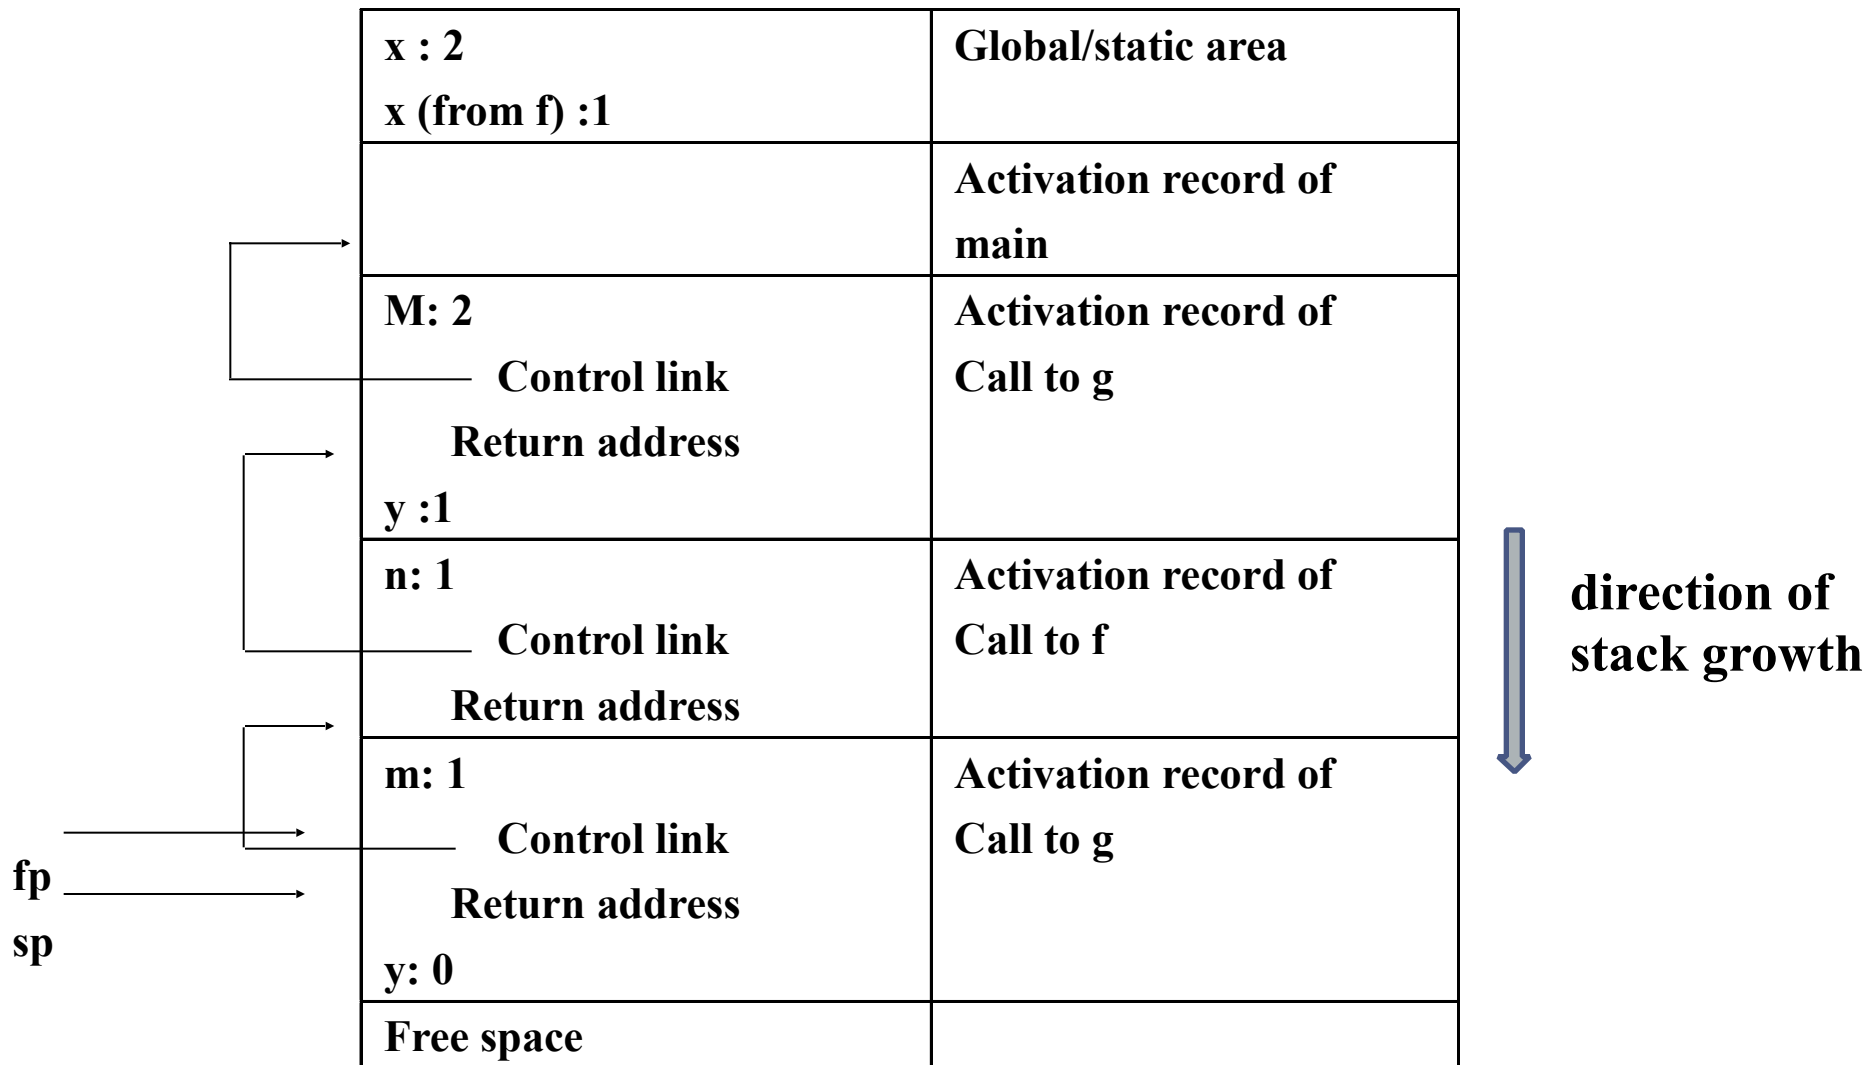
The environment during the third call:



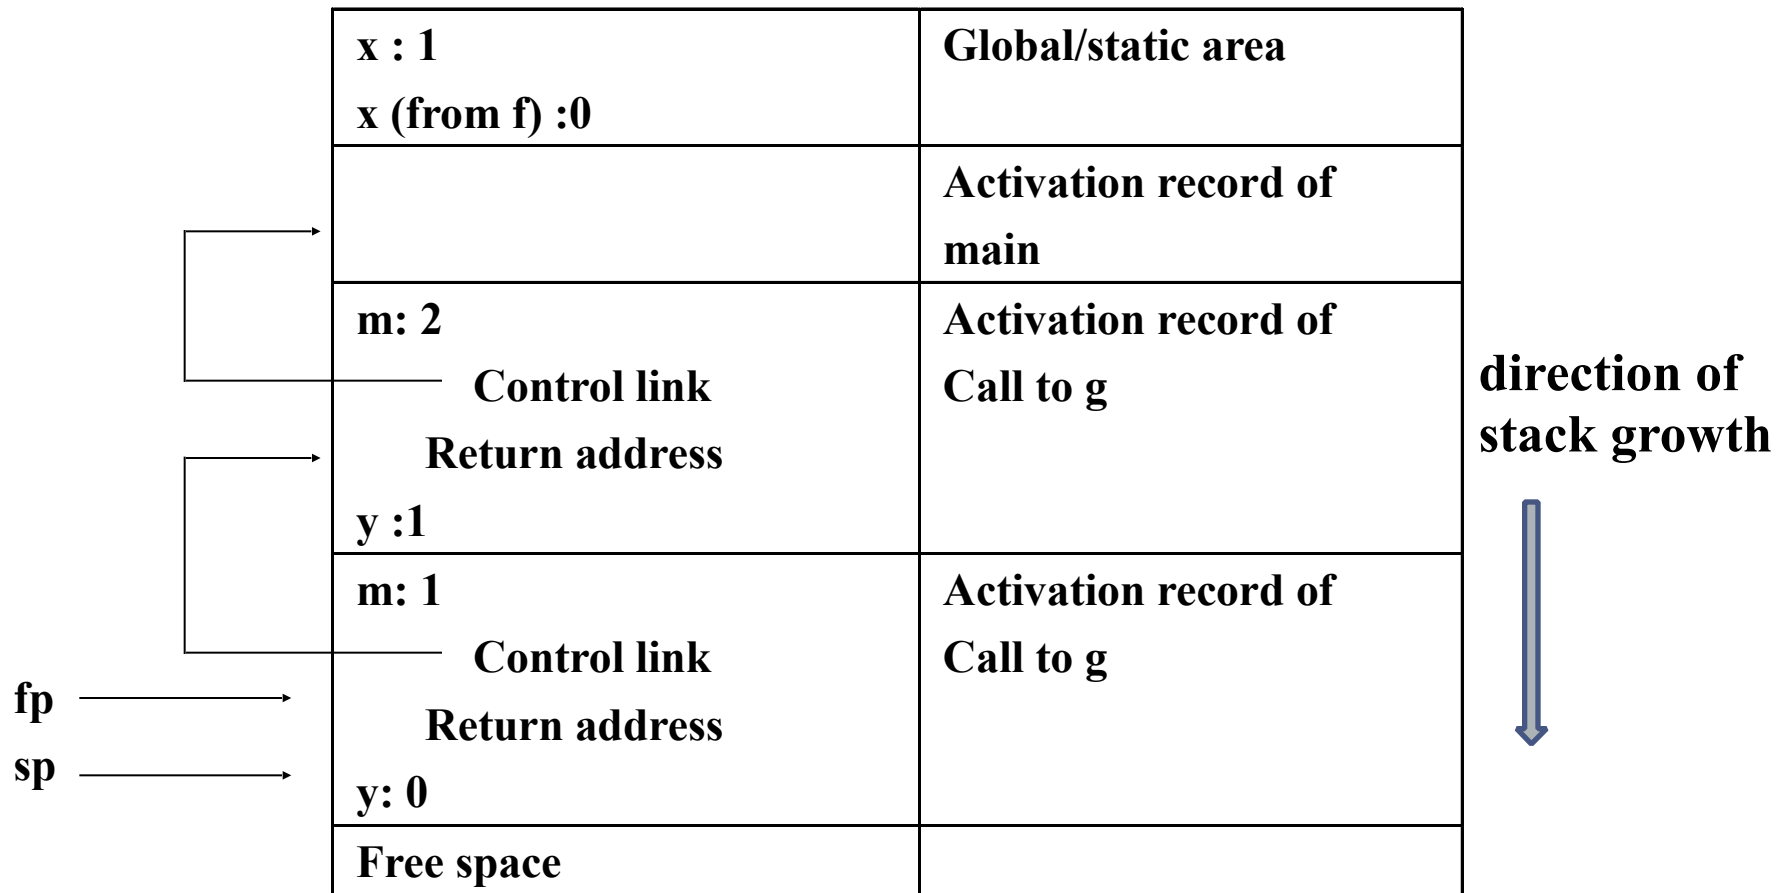
After the final call to gcd, each of the activations is removed in turn from the stack.

```
Example:  Int x=2;  
void g(int);/*prototype*/  
void f(int n)  
{  
    static int x=1;  
    g(n);  
    x--;  
}  
void g(int m)  
{  
    int y=m-1;  
    if (y>0)  
    {  
        f(y);  
        x--;  
        g(y);  
    }  
}  
main( )  
{  
    g(x);  
    return 0;  
}
```

(a) Runtime environment of the above program during the second call to g




(b) Runtime environment of the above program during the third call to g



Activations tree: a useful tool for the analysis of complex calling structures

Example: activation trees for the program of figures 7.3 and 7.5

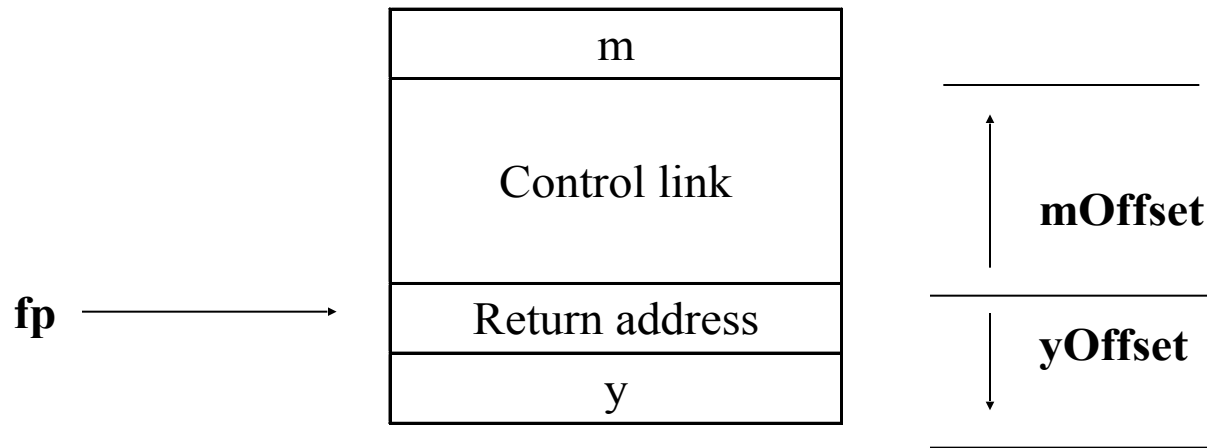


Note: In general, the stack of activation records at the beginning of a particular call has a structure equivalent to the path from the corresponding node in the activation tree to the root.

- Access to Names:
 - Parameters and local variable can **no longer be accessed by fixed addresses**
 - They must be found by offset from the current frame pointer.
 - In most language, **the offset can be statically computable** by the compiler.



Consider the procedure *g* in the C program of Figure 7.5.



Note:

- Each activation record of *g* has exactly the same form, and the parameter *m* and the local variable *y* are always in exactly the same relative location in the activation record.
- Both *m* and *y* can be accessed by their fixed offset from the *fp*.
- We have $mOffset = +4$ and $yOffset = -6$.
- The references to *m* and *y* can be written in machine code as $4(fp)$ and $-6(fp)$.

Example 7.4 Consider the C procedure

```
Void f(int x, char c)
```

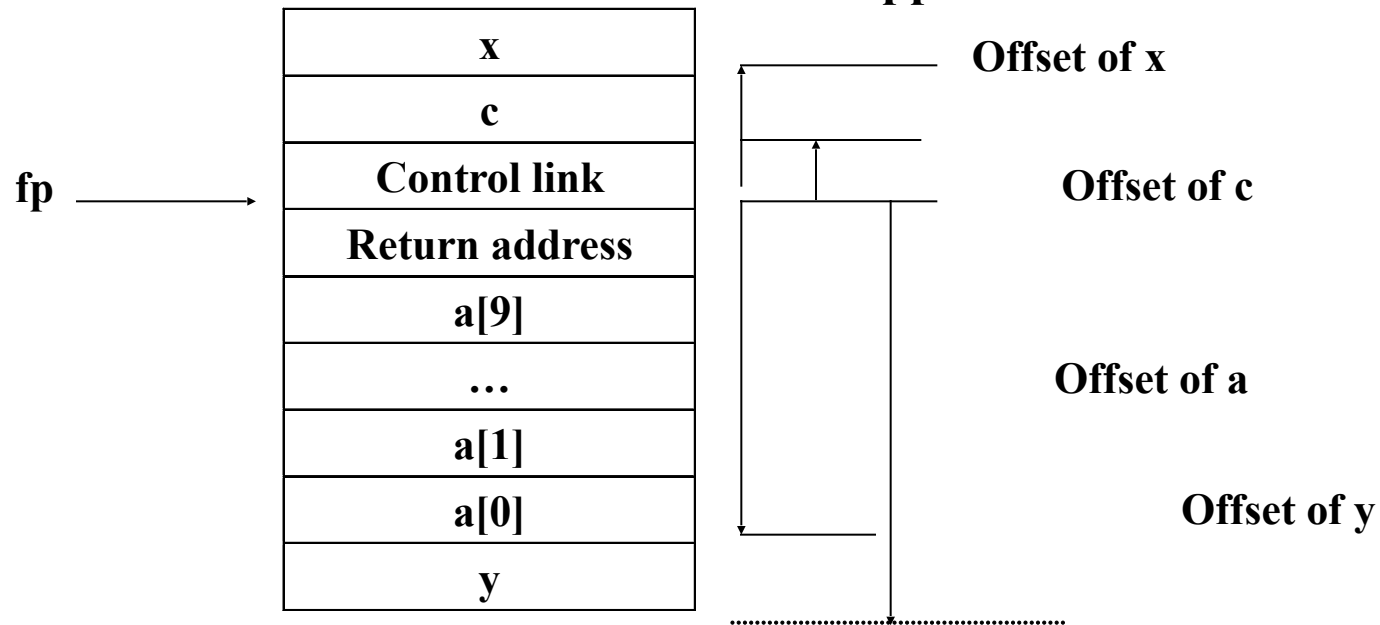
```
{ int a[10];
```

```
  double y;
```

```
  ...
```

```
}
```

The activation record for a call to f would appear as



Assuming two bytes for integers, four bytes for addresses, one byte for character and eight bytes for double-precision floating point, we would have the following offset values:

Name	Offset
x	+5
c	+4
a	-24
y	-32

Now, an access of $a[i]$, would require the computation of the address:

$$(-24 + 2 * i)(fp)$$

The calling sequence:

▫ When a procedure is called

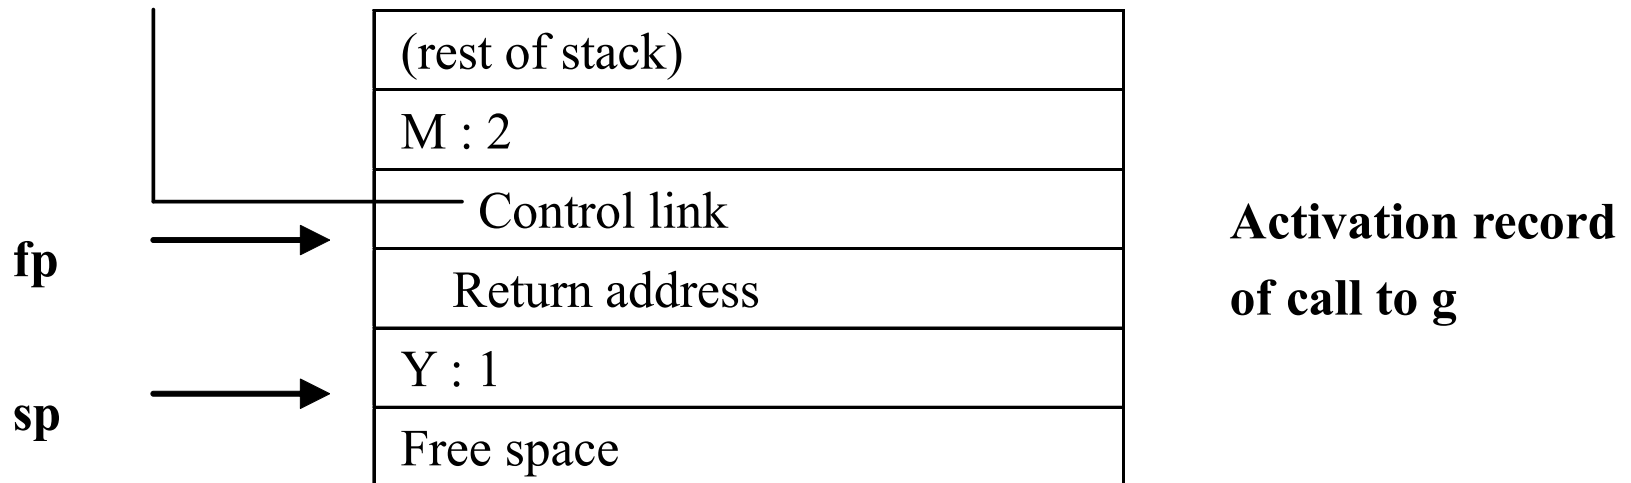
- **Compute the arguments** and store them in their correct positions in the new activation record of the procedure;
- **Store the fp** as the control link in the new activation record;
- **Change the fp** so that it points to the beginning of the new activation record;
- **Store the return address** in the new activation record;
- **Perform a jump** to the code of the procedure to be called.

▫ When a procedure exits

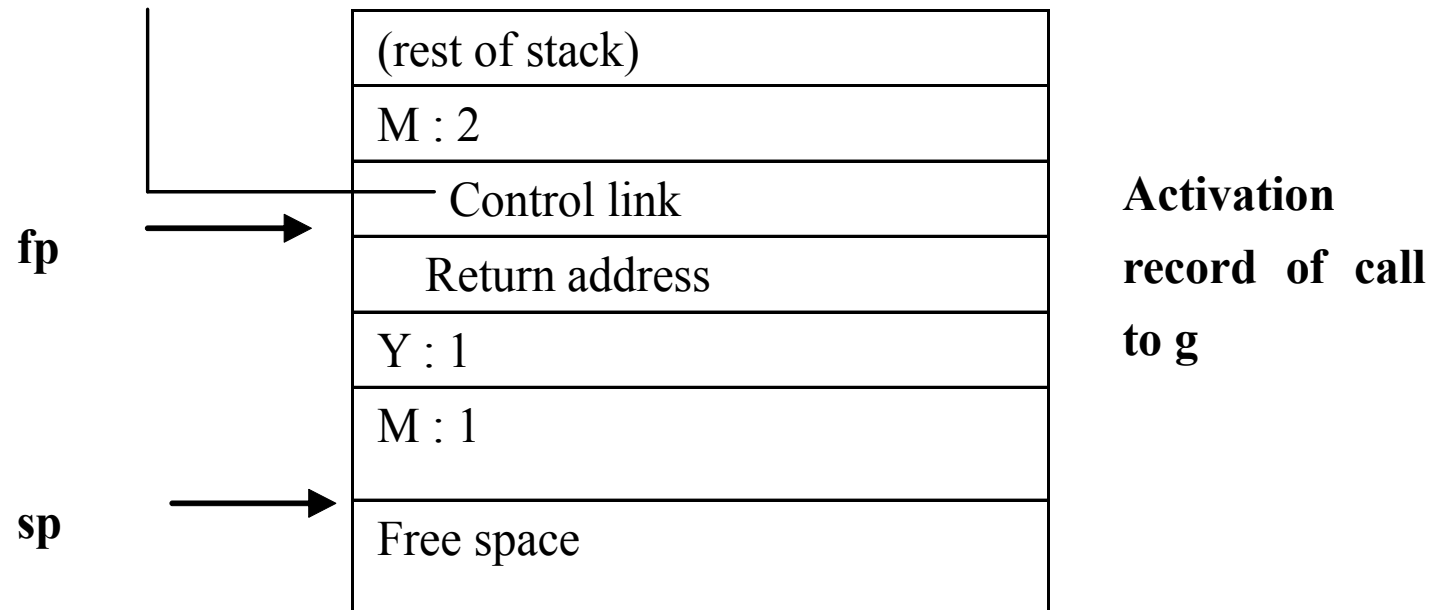
- Copy the fp to the sp.
- Load the control link into the fp.
- Perform a jump to the return address;
- Change the sp to pop the arguments.



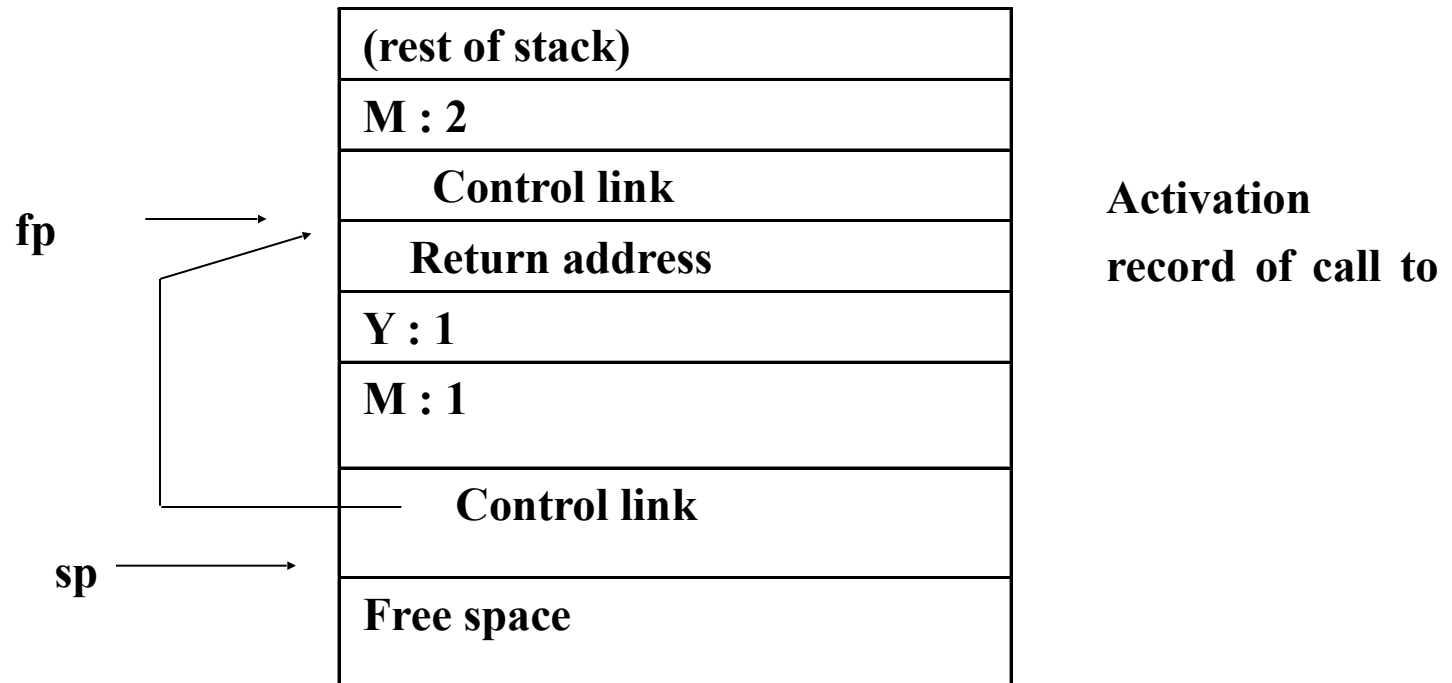
Example 7.5 Consider the situation just before the last call to g



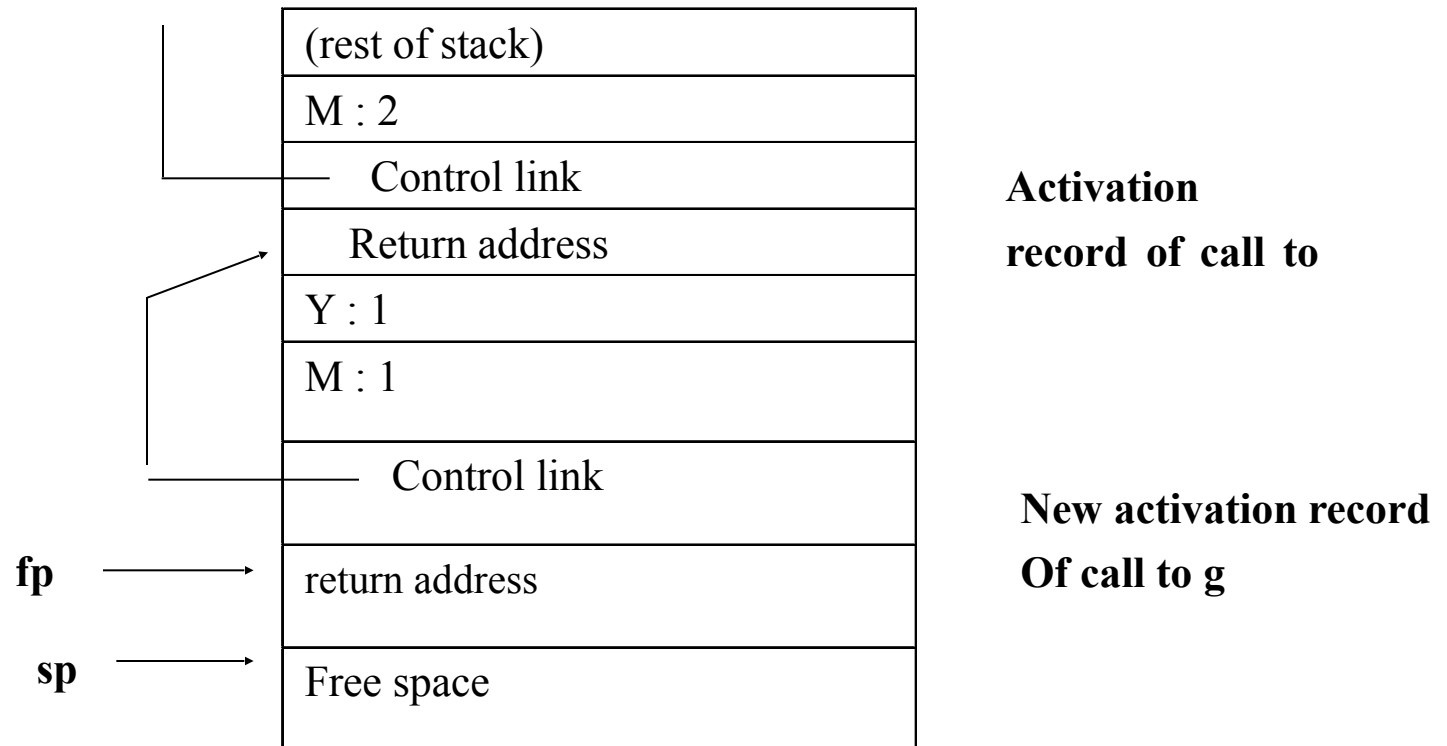
As the new call to g is made, first the value of parameter m is pushed onto the runtime stack:



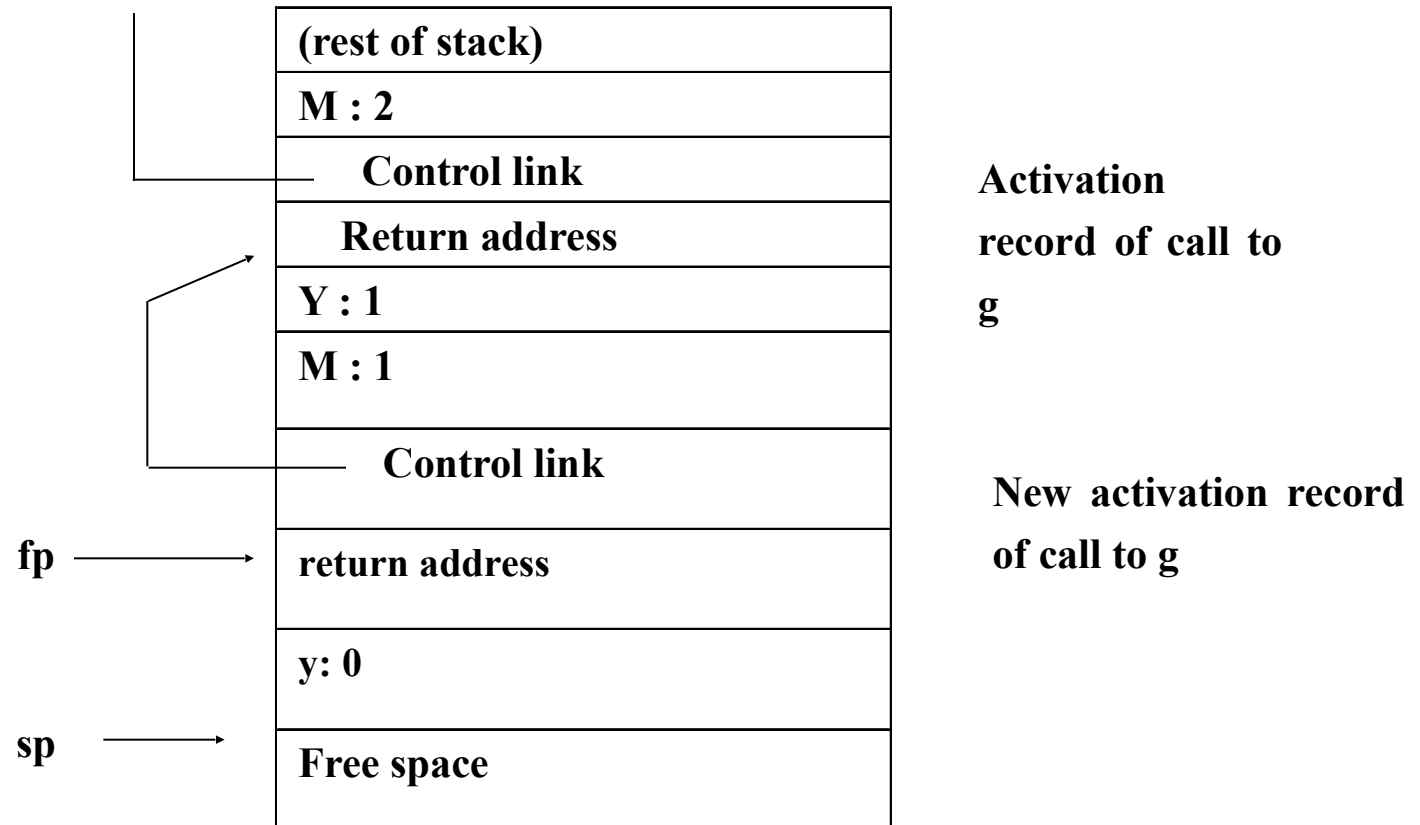
Then the fp is pushed onto the stack:



Now, the sp is copied into the fp, the return address is pushed onto the stack, and the jump to the new call of g is made:



Finally, g allocates and initializes the new y on the stack to complete the construction of the new activation record:



- **Dealing with variable-length data**

- The number of arguments in a call may vary from call to call, and
- The size of an array parameter or a local array variable may vary from call to call
- An example of situation 1 is the printf function in C:
 - `Printf(“%d%s%c”, n, prompt, ch)`
 - Has four arguments, while
 - `Printf(“Hello, world\n”)`
 - Has only one argument



- C compiler typically deal with this by **pushing the arguments to a call in reverse order** onto the runtime stack.
- The **first parameter is always located at a fixed offset from the fp** in the implementation described above
- Another option is to use a processor mechanism such as **ap (argument pointer)** in VAX architecture.



- An example of situation 2 is the unconstrained array of Ada:

Type int_vector is

Array(INTEGER range <>) of INTEGER;

Procedure sum (low, high: INTEGER;

A: Int_vector) return INTEGER

Is

Temp: Int_Array (low..high);

Begin

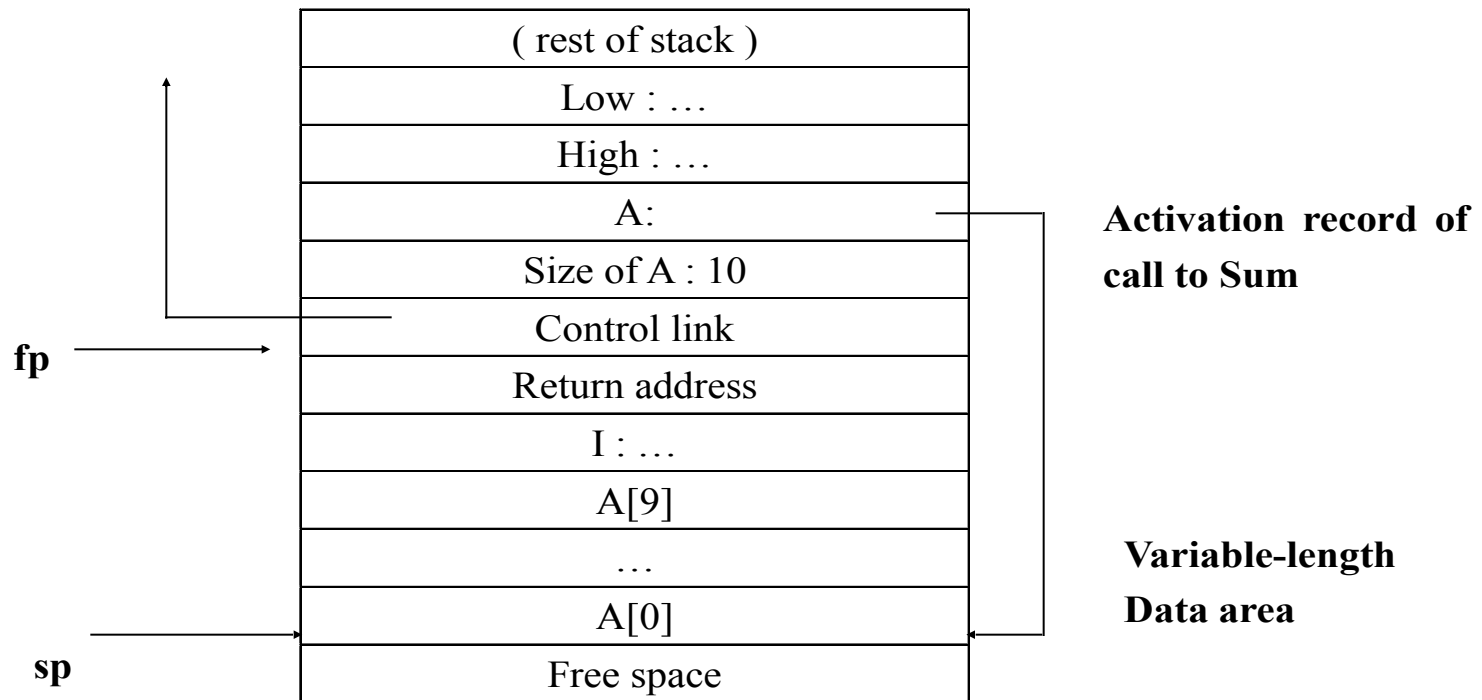
...

end sum;

- A typical method is to **use an extra level of indirection** for the **variable-length data, storing a pointer to the actual data** in a location that can be predicated at compile time.



We could implement an activation record for SUM as follows:



Now, for instance, access to $A[i]$ can be achieved by computing

$$@6(fp) + 2 * i$$

- Note:
 - In the implementation described in the previous example, the caller must **know the size** of any activation record of Sum
 - The size of the parameter part and the bookkeeping part is known to the compiler at the point of call
 - The size of the local variable part is not, in general, known at the point of call. Variable-length local variables can be dealt with in a similar way



- **Local Temporaries and Nested Declarations:**

- Two more complications to the basic stack-based runtime environment

(1) Local temporaries are partial results of computations that must be saved across procedure calls, for example:

$$x[i] = (i + j) * (i/k + f(j))$$

- The three partial results need to be saved across the call to f:

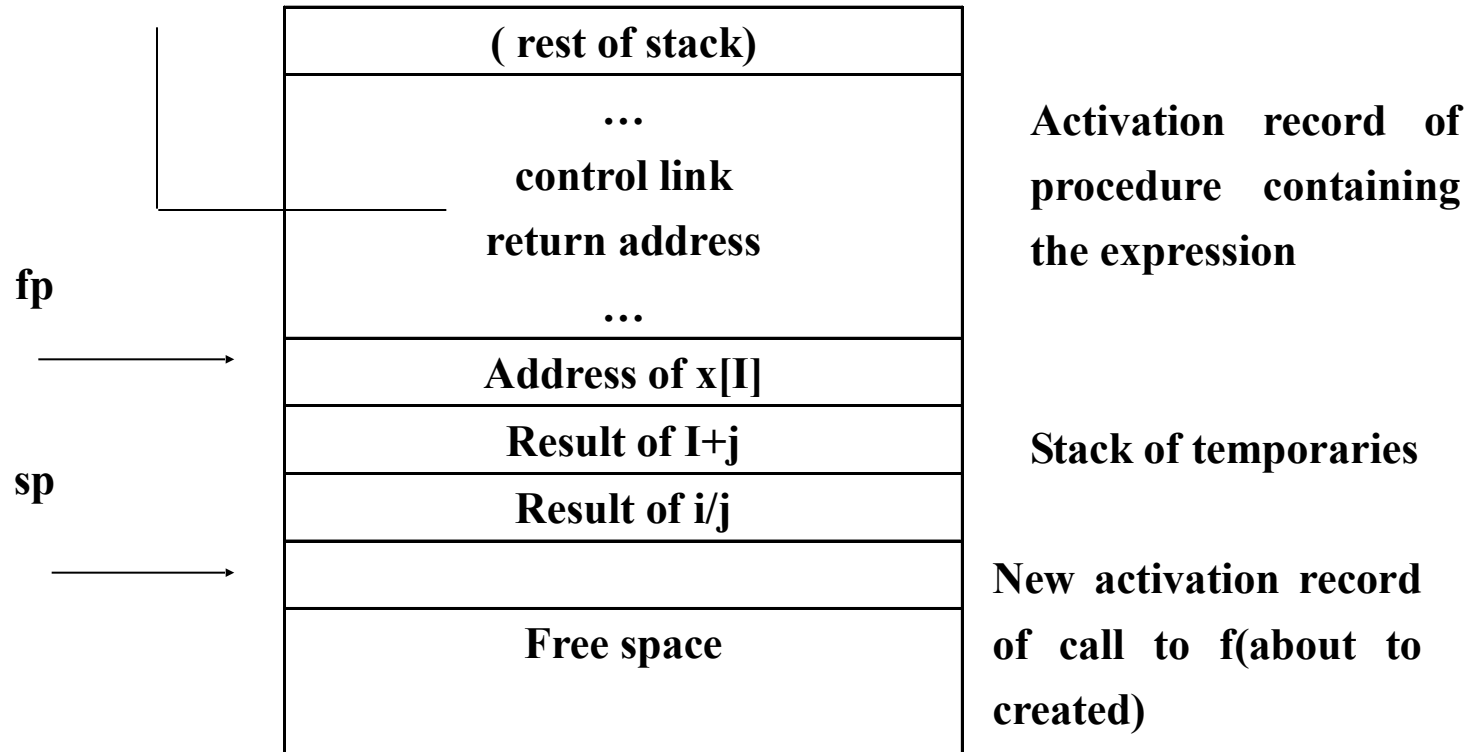
The address $x[i]$;

The sum $i+j$;

The quotient i/k ;



The runtime stack might appear as follows at the point just before the call to f:



- Nested declarations present a similar problem. Consider the C code

```

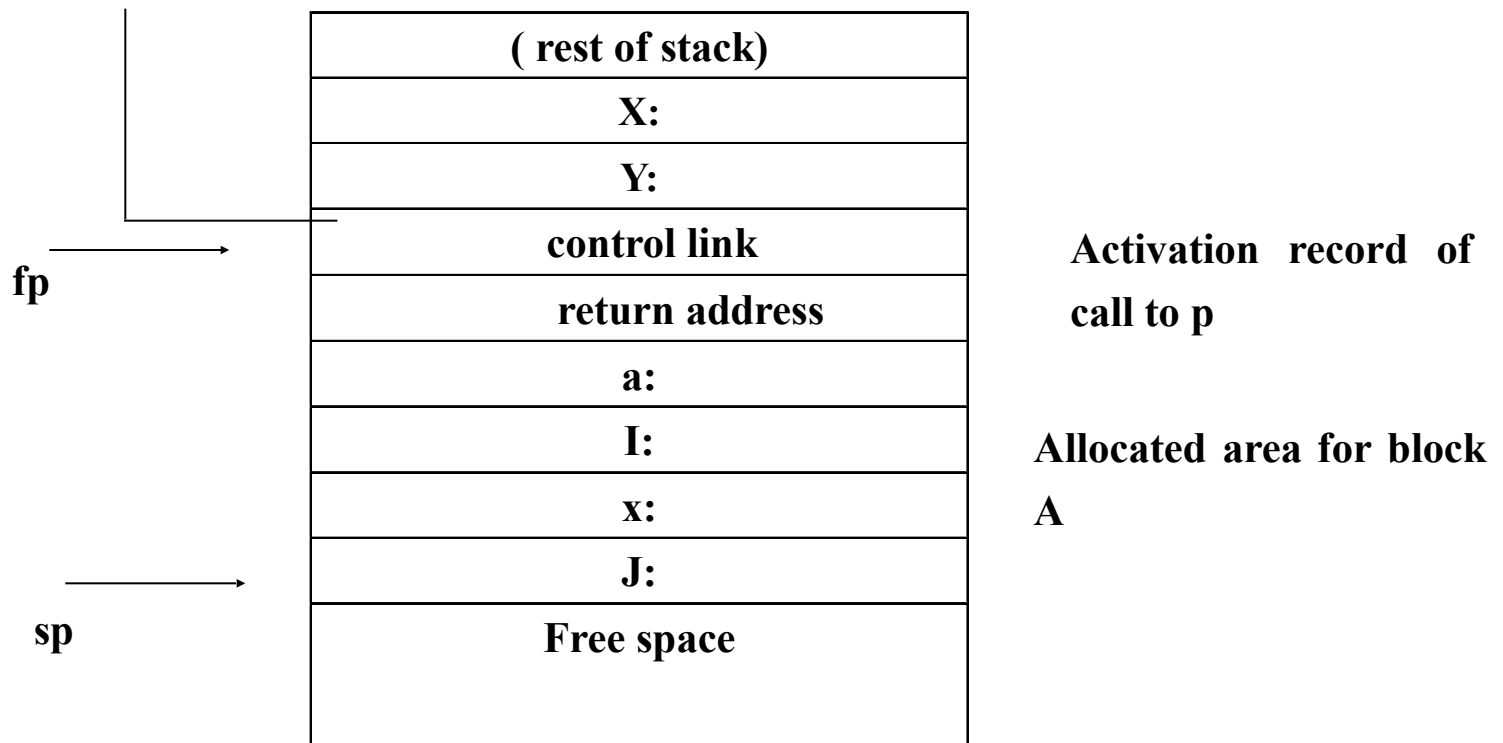
Void p( int x, double y)
{
    char a;
    int I;
    ...
    A:{    double x;
           Int j;
           ...
        }
    ...
    B:{
           char *a;
           Int k;
           ...
        }
    ...
}

```

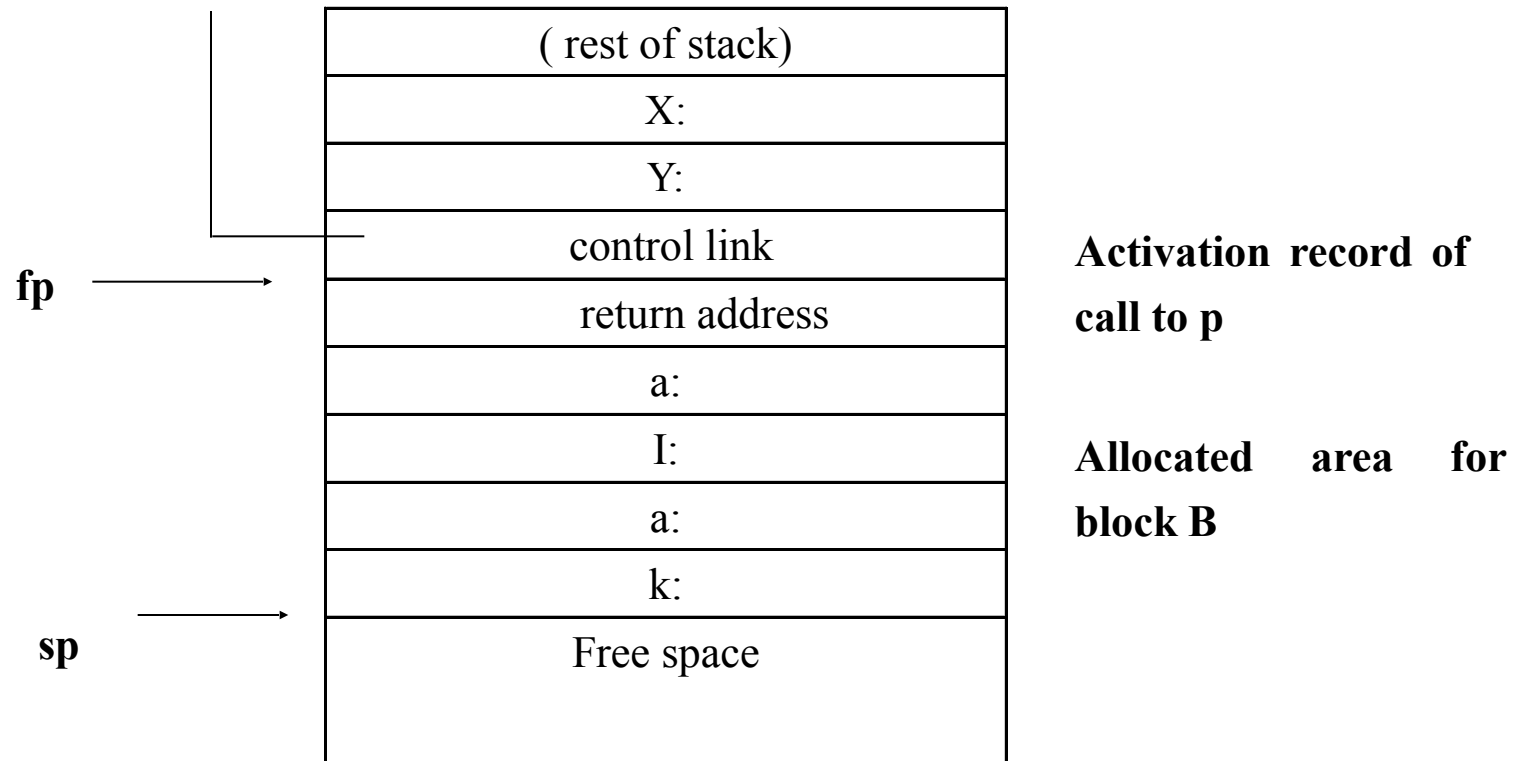
- There are two block (also called compound statement), labeled A and B nested inside the body of procedure p.
- The nested local declaration does not need to be allocated until entered;
- Both A and B do not need to be allocated simultaneously.



A simpler method is to treat them in a similar way to temporary expression. For instance, just after entering block A in the sample C just given, the runtime stack would appear as follows:



And just after entry to block B it would look as follows:



7.3.2 Stack-Based Environment with local Procedures

- Consider the non-local and non-global references
- Example: Pascal program showing nonlocal, nonglobal reference

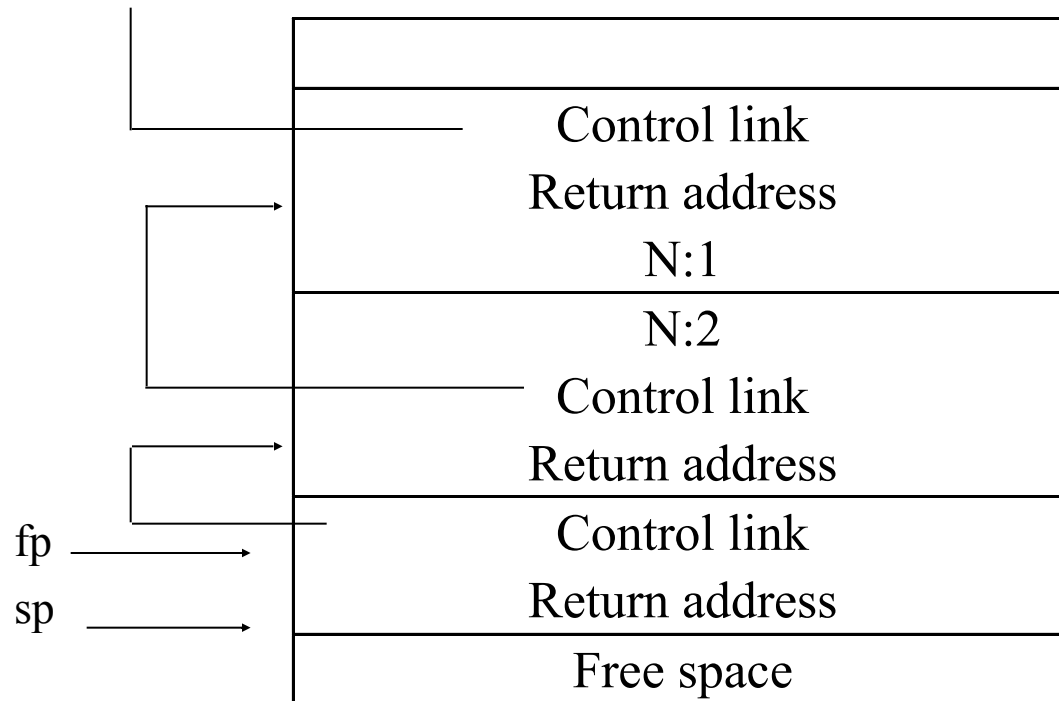
```

Program nonlocalRef;
Procedure P;
  Var N: integer;
    Procedure Q;
      Begin
        (* a reference to N is now non-local and non-global *)
      end; (* q *)
    Procedure R(N: integer);
      Begin
        Q;
      End; (* r *)
    Begin (* p *)
      N:=1;
      R(2);
    End; (* p *)
  Begin (* main *)
    P;
  End.

```



The runtime stack for the program above:



Activation record of
main program

Activation record of
Call to p

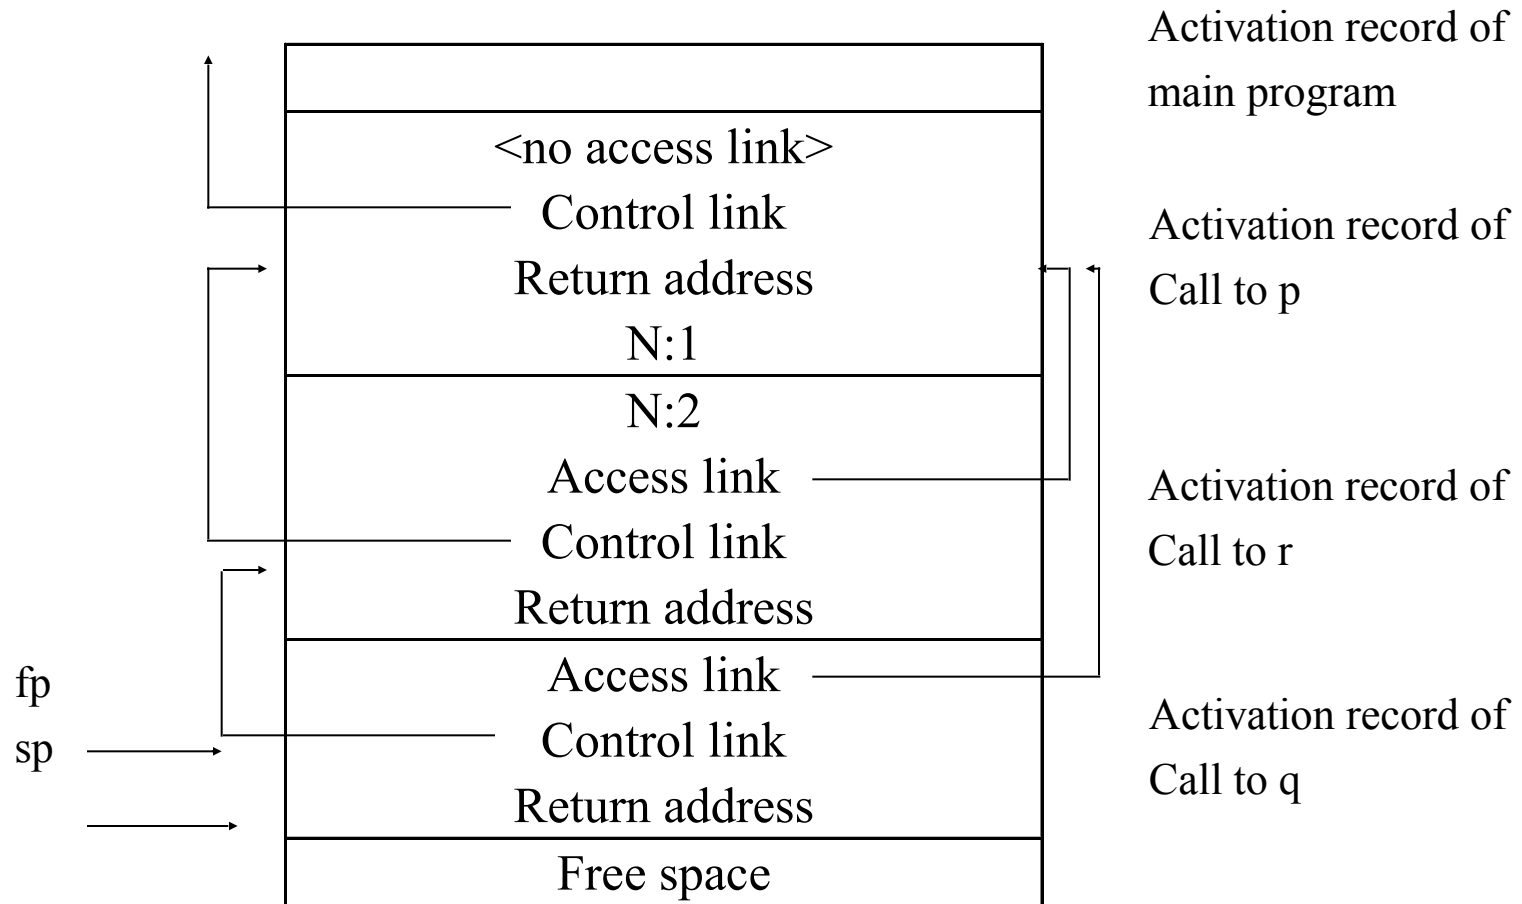
Activation record of
Call to r

Activation record of
Call to q

- **N cannot be found using any of the bookkeeping** information that is kept in the runtime environment up to now
- To solve the above problem about variable access, we add an extra piece of bookkeeping information called the ***access link*** to each activation record
- **Access link** represents the **defining environment** of the procedure;
- **Control link** represents the **calling environment** of the procedure.



The runtime stack for the program above with access links added.



- **Note:**

- The activation record of procedure p itself contains no access link, as any nonlocal reference with p **must be a global reference** and is accessed via the global reference mechanism
- This is the simplest situation, where the nonlocal reference is to a declaration **in the next outermost scope.**



Example: Pascal code demonstrating access chaining

Program chain

Procedure p;

Var x: integer;

 Procedure q;

 Procedure r;

 Begin

 X:=2;

 ...

 if ...then p;

 end;(*r*)

 begin

 r;

 end;(*q*)

begin

 q;

end;(*p*)

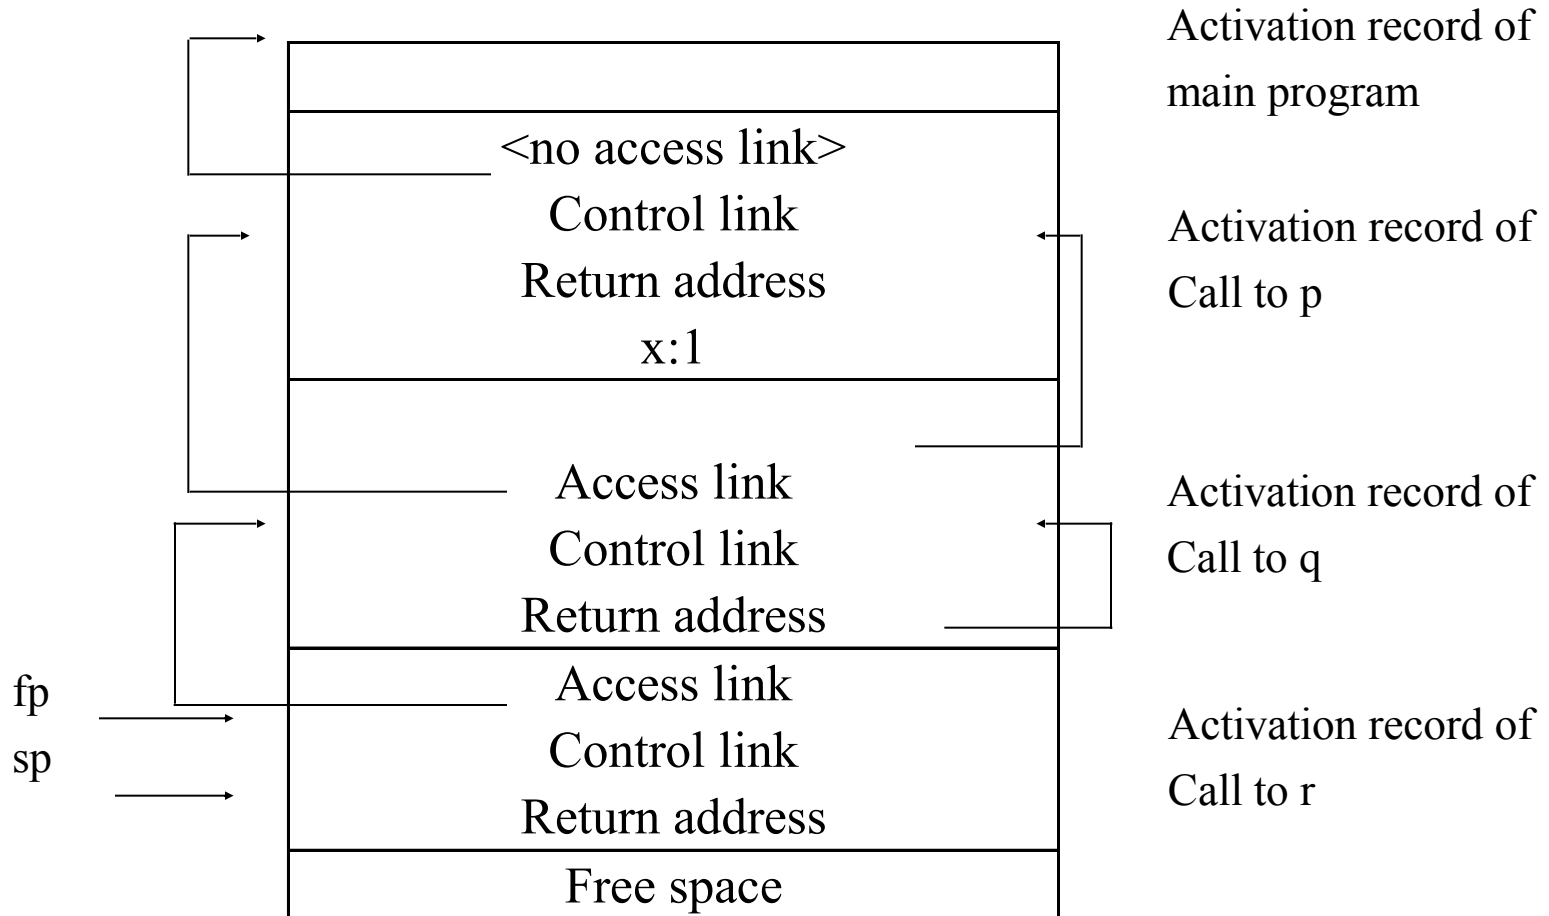


```
begin(* main *)  
    p;  
end.
```

- In this code, the assignment to x inside r, which refers to the x of p, must **traverse two scope levels** to find x
- In this environment, x must be reached by following two access links, a process that is called **access chaining**



The runtime stack after the first call to r:



- Note:

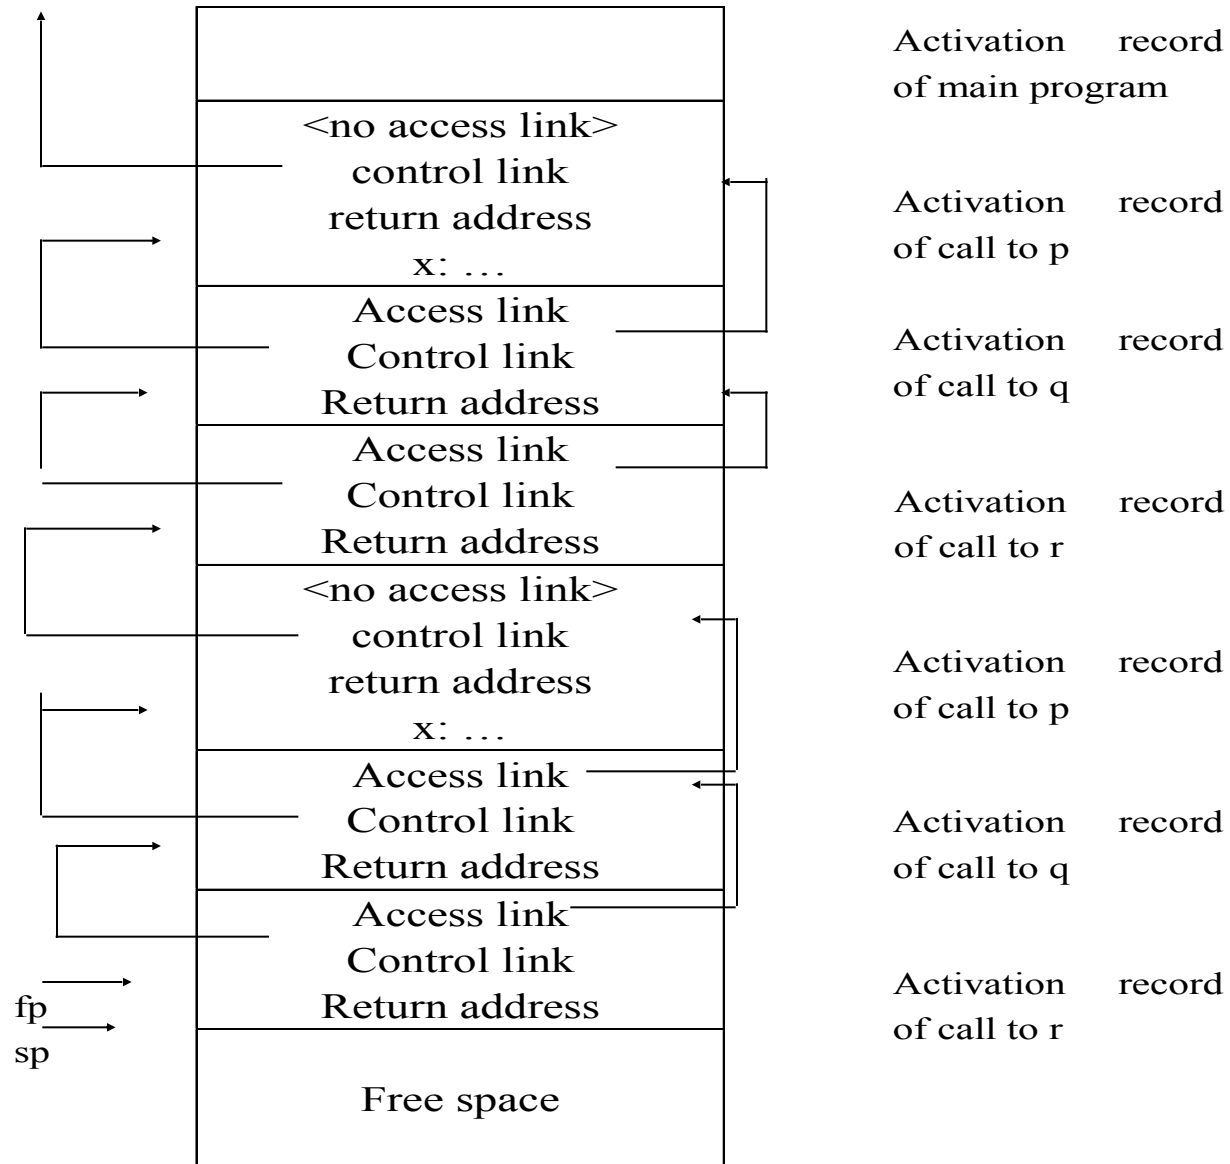
- The **amount of chaining, determined by** comparing the nesting level at the point of access with one of the declaration of the name
- In the above situation, the assignment to x is at nesting level 3, and x is declared at the nesting level 1, so two access links must be followed
- However, the access chaining is an inefficient method for variable access, for each nonlocal reference with a large nesting difference, a lengthy sequence of instruction must be executed.
- There is a method of implementing **access links in a lookup table indexed by nesting level**, called **display**, to avoid the execution overhead of chaining.



- **The calling sequence**
- The **changes needed** to implement access links:
 - (1) The access link must be pushed onto the runtime stack just before the fp during a call
 - (2) The sp must be adjusted by an extra amount to remove the access link after an exit
- How to find the access link of a procedure during a call?
 - Using the (compile-time) nesting level information attached to the declaration of the procedure
 - Generate an access chain as if to access a variable at the same nesting level
 - **The access link and the control link are the same, if the procedure is local**



Given the code of **PROGRAM CHAIN** the runtime stack after the second call to r (assuming a recursive call to p) would look as follows:



7.3.3 Stack-Based Environment with Procedure Parameters



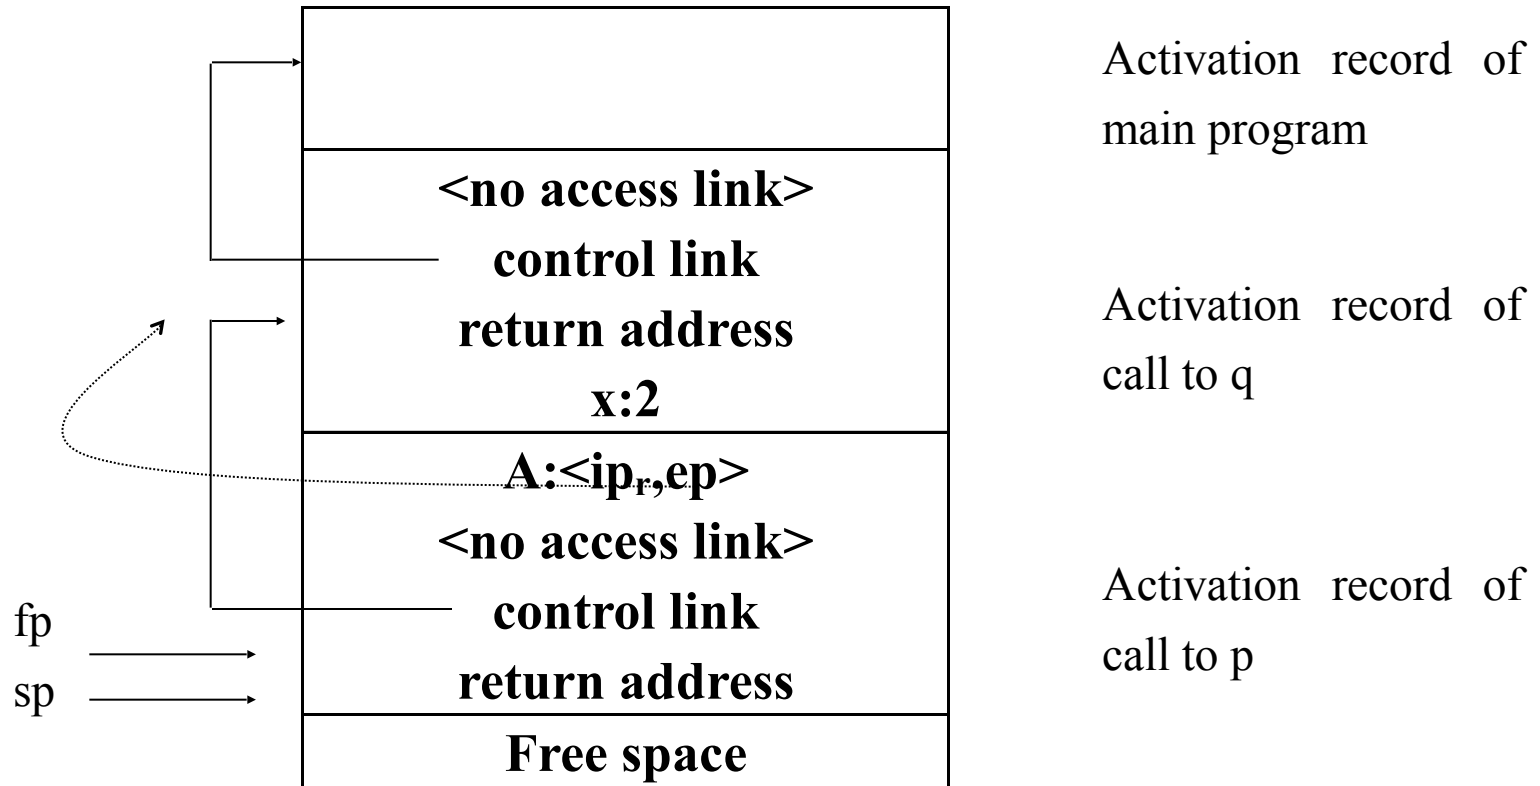
Example 7.7 Consider the standard pascal program of Figure 7.14, which has a procedure p, with a parameter *a* that is also a procedure.

```
Program closureEx(output)
  Procedure p (procedure a);
  Begin
    a;
  end;
  procedure q;
  var x:integer;
    procedure r;
    begin
      writeln(x);
    end;
  begin
    x:=2;
    p(r);
  end; (* q *)
  begin (* main *)
    q;
  end.
```

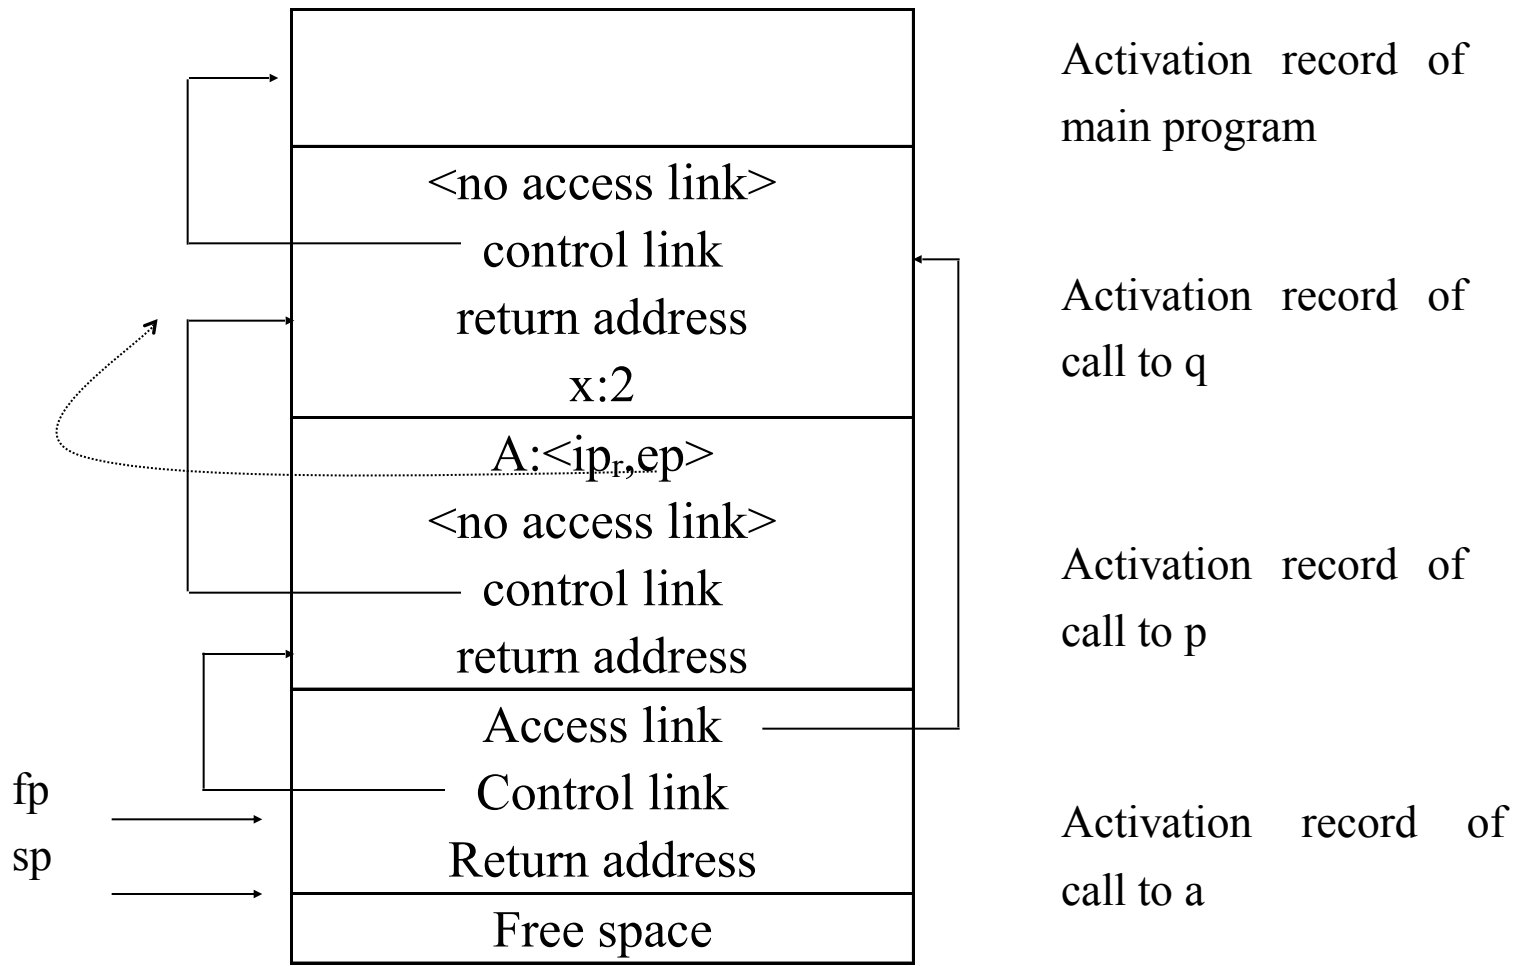


...

The runtime stack just after the call to p in the code of Figure 7.14



... **The runtime stack just after the call to a in the code of Figure 7.14**

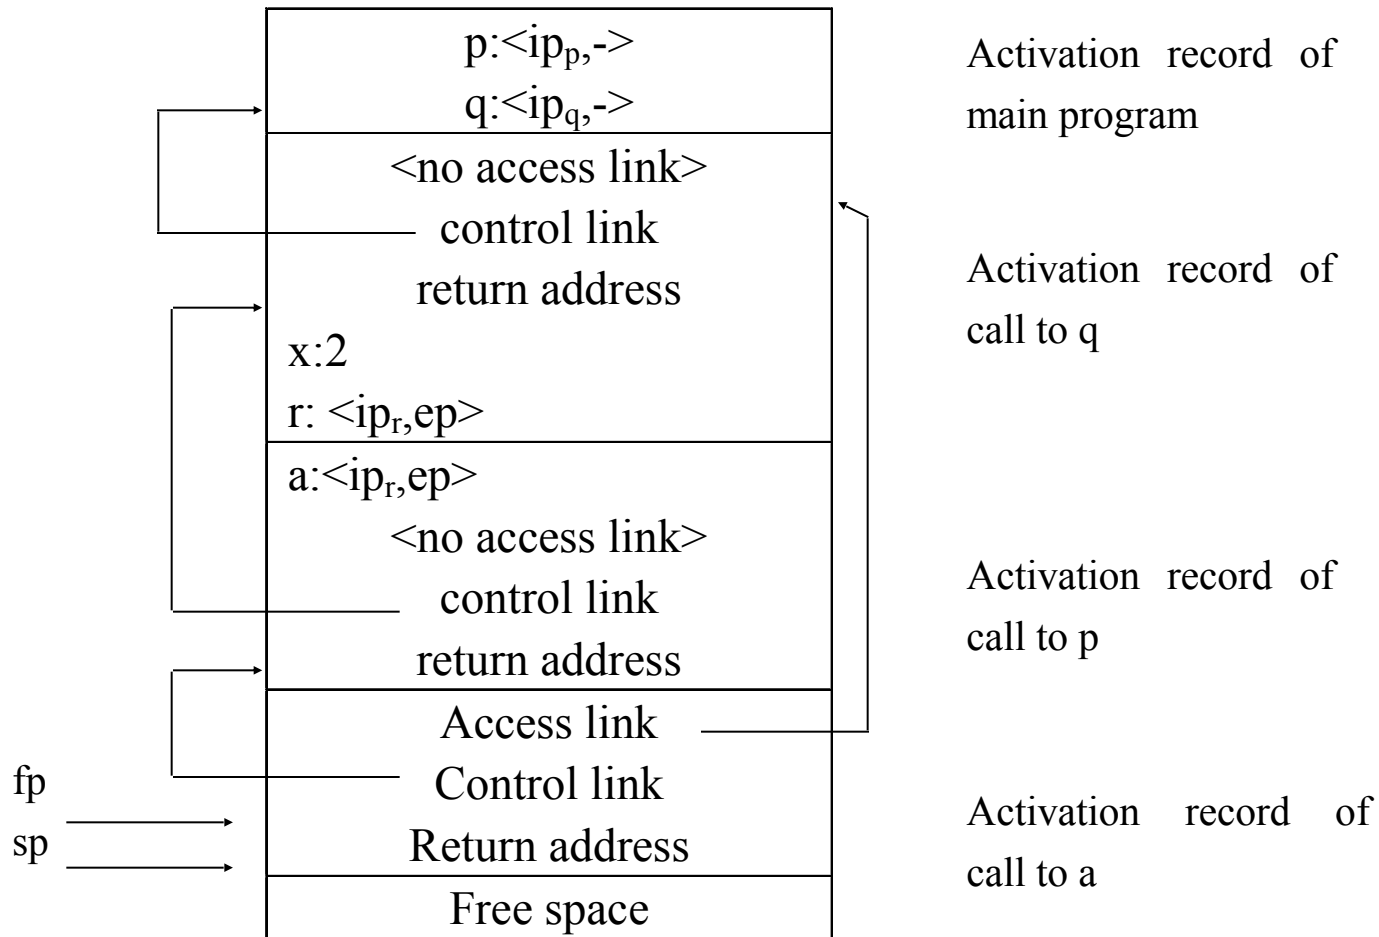


- Note:

- The calling sequence must distinguish clearly between ordinary procedures and procedure parameters
- When calling ordinary procedure, fetch the access link using the nesting level and jump directly to the code of the procedure
- A procedure parameter has **its access link stored in the local activation record**, and an indirect call must be performed to the ip stored in the current activation record
- If all procedure values are stored in the environment as closures, the following page shows the environment



Runtime stack just after the call to a in the code of Figure 7.14 with all procedures kept as closures in the environment



End of Part One

THANKS